

Notes on Probabilistic Machine Learning

Probability theory is nothing but common sense reduced to calculations.

— Pierre Laplace, 1812

Contents

I	Mathematical Foundations	3
1	Multivariate Probability Models	3
1.1	Measures the dependency	3
1.2	The multivariate Gaussian/Normal distribution (MVN)	4
1.3	Marginals and conditionals of an MVN	4
1.4	Linear Gaussian systems and Bayes rule for Gaussians	5
2	Maximum Likelihood Estimation	7
3	Matrix Calculus	8
3.1	Functions that map scalars to scalars	8
3.2	Functions that map vectors to scalars	8
3.3	Functions that map vectors to vectors	9
3.4	Functions that map matrices to scalars	9
4	Bound Optimization	11
4.1	EM algorithm	11
5	Mixture Models	13
5.1	Gaussian mixture models	13
5.2	EM for a GMM	13
5.3	Bernoulli mixture model	14
II	Linear Models and Neural Networks	15
6	Generalized Linear Models	15
6.1	Logistic regression	16
6.2	Linear regression	16
6.3	Binomial regression	16
6.4	Poisson regression	17
6.5	Canonical and non-canonical link functions	18
6.6	Maximum likelihood estimation	18

III	Dimentionality Reduction	20
7	Principal Component Analysis	20
7.1	PCA setup	20
7.2	Algorithm derivation	21
7.3	Computational tricks	22
8	Factor Analysis	24
8.1	Prerequisite: evidence lower bound (ELBO)	24
8.2	EM algorithm for FA	25
8.3	EM algorithm for PPCA	26
8.4	Limiting case: EM algorithm for PCA	26
8.5	Special case: Probabilistic PCA	27
8.6	Summary	29
9	Autoencoders	30
9.1	Deterministic autoencoders	30
9.2	Variational autoencoders (VAE)	32
IV	Kernel Methods	36
10	Mercer Kernels and Mercer's Theorem	36
10.1	Mercer's theorem	36
10.2	Mercer kernels examples	37
10.3	Combining kernels	37
11	Gaussian Processes	38
11.1	Noise-free observations	38
11.2	Noisy observations	39
11.3	Bayesian linear regression	39
11.4	Estimating the kernel	40

Part I

Mathematical Foundations

This part is not aiming to provide a general coverage of all the mathematical foundations, but rather to provide a quick reference for the mathematical concepts that are frequently used and worth noting.

1 Multivariate Probability Models

1.1 Measures the dependency

Covariance

The covariance between two rv's X and Y measures the degree to which X and Y are linearly related.

$$\text{Cov}(X, Y) \triangleq \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])] = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$$

If $\mathbf{x} \in \mathbb{R}^D$ is a random vector, its covariance matrix is defined to be the following symmetric, positive semi-definite matrix:

$$\begin{aligned} \text{Cov}[\mathbf{x}] &\triangleq \mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^\top] \triangleq \mathbf{\Sigma} = \mathbf{K}_{xx} \\ &= \begin{pmatrix} \mathbb{V}[X_1] & \text{Cov}[X_1, X_2] & \cdots & \text{Cov}[X_1, X_D] \\ \text{Cov}[X_2, X_1] & \mathbb{V}[X_2] & \cdots & \text{Cov}[X_2, X_D] \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}[X_D, X_1] & \text{Cov}[X_D, X_2] & \cdots & \mathbb{V}[X_D] \end{pmatrix} \end{aligned}$$

expanding $\mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^\top]$ gives the important result

$$\mathbf{R}_{xx} = \mathbb{E}[\mathbf{x}\mathbf{x}^\top] = \mathbf{\Sigma} + \boldsymbol{\mu}\boldsymbol{\mu}^\top$$

Under linear transformations, the covariance matrix transforms as follows:

$$\text{Cov}[\mathbf{A}\mathbf{x} + \mathbf{b}] = \mathbf{A}\text{Cov}[\mathbf{x}]\mathbf{A}^\top$$

The cross-covariance between two random vectors

$$\text{Cov}[\mathbf{x}, \mathbf{y}] = \mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{y} - \mathbb{E}[\mathbf{y}])^\top] = \text{Cov}[\mathbf{x}, \mathbf{y}]^\top$$

where $\text{Cov}[\mathbf{x}, \mathbf{y}] \in \mathbb{R}^{\text{Dim}(\mathbf{x}) \times \text{Dim}(\mathbf{y})}$ and the (i, j) -th entry is $\text{Cov}[\mathbf{x}, \mathbf{y}]_{ij} = \text{Cov}[x_i, y_j]$

Correlation

The correlation coefficient between X and Y is defined as

$$\rho \triangleq \text{corr}[X, Y] = \frac{\text{Cov}[X, Y]}{\sqrt{\mathbb{V}[X]\mathbb{V}[Y]}}$$

which is a normalised measure with $-1 \leq \rho \leq 1$. This can be thought as a measure of *degree of linearity*. (i.e. $\text{corr}[X, Y] = 1$ if and only if $Y = aX + b$ (and $a > 0$) for some parameters a and b)

Correlation matrix of a vector $\mathbf{x}$ of D related random variables is defined as

$$\text{corr}(\mathbf{x}) = \begin{pmatrix} 1 & \rho_{12} & \cdots & \rho_{1D} \\ \rho_{21} & 1 & \cdots & \rho_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ \rho_{D1} & \rho_{D2} & \cdots & 1 \end{pmatrix}$$

or compactly as,

$$\text{corr}(\mathbf{x}) = \text{diag}(\mathbf{K}_{xx})^{-\frac{1}{2}} \mathbf{K}_{xx} \text{diag}(\mathbf{K}_{xx})^{-\frac{1}{2}}$$

where $\mathbf{K}_{xx}$ is the (auto)covariance matrix of $\mathbf{x}$,

$$\begin{aligned} \mathbf{K}_{xx} &= \mathbf{\Sigma} = \mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^\top] \\ &= \mathbb{E}[\mathbf{x}\mathbf{x}^\top] - \boldsymbol{\mu}\boldsymbol{\mu}^\top = \mathbf{R}_{xx} - \boldsymbol{\mu}\boldsymbol{\mu}^\top \end{aligned}$$

where $\mathbf{R}_{xx}$ is the (auto)correlation matrix of $\mathbf{x}$.

Uncorrelated does not imply independent

Independent implies uncorrelated,

$$p(X, Y) = p(X)p(Y) \implies \text{Cov}[X, Y] = 0 \implies \text{corr}[X, Y] = 0$$

However, the converse is not true: uncorrelated does not imply independent (eg. $X \sim U(-1, 1)$, $Y = X^2$ have zero correlation but are not independent). A more general measure of dependence between random variables is mutual information. This is zero only if the variables truly are independent.

1.2 The multivariate Gaussian/Normal distribution (MVN)

The MVN density is defined by the following:

$$\mathcal{N}(\mathbf{y}|\boldsymbol{\mu}, \mathbf{\Sigma}) \triangleq \frac{1}{(2\pi)^{D/2}|\mathbf{\Sigma}|^{1/2}} \exp\left[-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})^\top \mathbf{\Sigma}^{-1}(\mathbf{y} - \boldsymbol{\mu})\right]$$

where $\boldsymbol{\mu} = \mathbb{E}[\mathbf{y}] \in \mathbb{R}^D$ is the mean vector, and $\mathbf{\Sigma} = \text{Cov}[\mathbf{y}]$ is the $D \times D$ covariance matrix.

1.3 Marginals and conditionals of an MVN

For a Gaussian random vector $\mathbf{y} = (\mathbf{y}_1, \mathbf{y}_2)$, if we are conditioned on a set of noise-free observations $\mathbf{y}_2$, then we can infer the hidden parts $\mathbf{y}_1$ by the posterior $p(\mathbf{y}_1|\mathbf{y}_2)$.

Suppose $\mathbf{y} = (\mathbf{y}_1, \mathbf{y}_2)$ is jointly Gaussian with parameters

$$\boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_1 \\ \boldsymbol{\mu}_2 \end{pmatrix}, \quad \mathbf{\Sigma} = \begin{pmatrix} \mathbf{\Sigma}_{11} & \mathbf{\Sigma}_{12} \\ \mathbf{\Sigma}_{21} & \mathbf{\Sigma}_{22} \end{pmatrix}, \quad \boldsymbol{\Lambda} = \mathbf{\Sigma}^{-1} = \begin{pmatrix} \boldsymbol{\Lambda}_{11} & \boldsymbol{\Lambda}_{12} \\ \boldsymbol{\Lambda}_{21} & \boldsymbol{\Lambda}_{22} \end{pmatrix}$$

where $\boldsymbol{\Lambda}$ is the **precision matrix** and $\mathbf{\Sigma}_{12} = \mathbf{\Sigma}_{21}^\top = \text{Cov}[\mathbf{y}_1, \mathbf{y}_2]$.

Then the marginals are given by

$$\boxed{\begin{aligned} p(\mathbf{y}_1) &= \mathcal{N}(\mathbf{y}_1|\boldsymbol{\mu}_1, \mathbf{\Sigma}_{11}) \\ p(\mathbf{y}_2) &= \mathcal{N}(\mathbf{y}_2|\boldsymbol{\mu}_2, \mathbf{\Sigma}_{22}) \end{aligned}}$$

and the posterior conditional is given by

$$\begin{aligned}
p(\mathbf{y}_1|\mathbf{y}_2) &= \mathcal{N}(\mathbf{y}_1|\boldsymbol{\mu}_{1|2}, \boldsymbol{\Sigma}_{1|2}) \\
\boldsymbol{\mu}_{1|2} &= \boldsymbol{\mu}_1 + \boldsymbol{\Sigma}_{12}\boldsymbol{\Sigma}_{22}^{-1}(\mathbf{y}_2 - \boldsymbol{\mu}_2) \\
&= \boldsymbol{\mu}_1 - \boldsymbol{\Lambda}_{11}^{-1}\boldsymbol{\Lambda}_{12}(\mathbf{y}_2 - \boldsymbol{\mu}_2) \\
&= \boldsymbol{\Sigma}_{1|2}(\boldsymbol{\Lambda}_{11}\boldsymbol{\mu}_1 - \boldsymbol{\Lambda}_{12}(\mathbf{y}_2 - \boldsymbol{\mu}_2)) \\
\boldsymbol{\Sigma}_{1|2} &= \boldsymbol{\Sigma}_{11} - \boldsymbol{\Sigma}_{12}\boldsymbol{\Sigma}_{22}^{-1}\boldsymbol{\Sigma}_{21} = \boldsymbol{\Lambda}_{11}^{-1}
\end{aligned}$$

The proofs rely on computing the Schur complement $\boldsymbol{\Sigma}/\boldsymbol{\Sigma}_{22} = \boldsymbol{\Sigma}_{11} - \boldsymbol{\Sigma}_{12}\boldsymbol{\Sigma}_{22}^{-1}\boldsymbol{\Sigma}_{21}$, which is not shown here. For the marginals, just extracting the rows and columns corresponding to $\mathbf{y}_1$ or $\mathbf{y}_2$. For the conditional: the conditional mean is a linear function of $\mathbf{y}_2$, and the conditional covariance is a constant matrix that is independent of $\mathbf{y}_2$.

1.4 Linear Gaussian systems and Bayes rule for Gaussians

Let $\mathbf{z} \in \mathbb{R}^L$ be an unknown vector of values, and $\mathbf{y} \in \mathbb{R}^D$ be some noisy measurement of $\mathbf{z}$. We assume these variables are related by the following joint distribution:

$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_z)$$

$$p(\mathbf{y}|\mathbf{z}) = \mathcal{N}(\mathbf{y}|\mathbf{W}\mathbf{z} + \mathbf{b}, \boldsymbol{\Sigma}_y)$$

where $\mathbf{W}$ is a matrix of size $D \times L$. This is an example of a **linear Gaussian system**.

The corresponding joint distribution, $p(\mathbf{z}, \mathbf{y}) = p(\mathbf{z})p(\mathbf{y}|\mathbf{z})$, is a $L + D$ dimensional Gaussian, with mean and covariance given by

$$\begin{aligned}
\boldsymbol{\mu} &= \begin{pmatrix} \boldsymbol{\mu}_z \\ \mathbf{W}\boldsymbol{\mu}_z + \mathbf{b} \end{pmatrix} \\
\boldsymbol{\Sigma} &= \begin{pmatrix} \boldsymbol{\Sigma}_z & \boldsymbol{\Sigma}_z\mathbf{W}^\top \\ \mathbf{W}\boldsymbol{\Sigma}_z & \boldsymbol{\Sigma}_y + \mathbf{W}\boldsymbol{\Sigma}_z\mathbf{W}^\top \end{pmatrix}
\end{aligned}$$

By applying the previous Gaussian conditioning formula to the joint $p(\mathbf{y}, \mathbf{z})$ we can compute the posterior $p(\mathbf{z}|\mathbf{y})$ as

$$\begin{aligned}
p(\mathbf{z}|\mathbf{y}) &= \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_{z|y}, \boldsymbol{\Sigma}_{z|y}) \\
\boldsymbol{\mu}_{z|y} &= \boldsymbol{\mu}_z + \boldsymbol{\Sigma}_z\mathbf{W}^\top(\boldsymbol{\Sigma}_y + \mathbf{W}\boldsymbol{\Sigma}_z\mathbf{W}^\top)^{-1}(\mathbf{y} - \mathbf{W}\boldsymbol{\mu}_z - \mathbf{b}) \\
\boldsymbol{\Sigma}_{z|y} &= \boldsymbol{\Sigma}_z - \boldsymbol{\Sigma}_z\mathbf{W}^\top(\boldsymbol{\Sigma}_y + \mathbf{W}\boldsymbol{\Sigma}_z\mathbf{W}^\top)^{-1}\mathbf{W}\boldsymbol{\Sigma}_z
\end{aligned}$$

However, this requires inverting a $D \times D$ matrix $(\boldsymbol{\Sigma}_y + \mathbf{W}\boldsymbol{\Sigma}_z\mathbf{W}^\top)$, which can be expensive if D is large. Instead, we can use the matrix inversion lemma (Woodbury matrix identity) to rewrite the posterior in a more efficient form:

$$\begin{aligned}
p(\mathbf{z}|\mathbf{y}) &= \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_{z|y}, \boldsymbol{\Sigma}_{z|y}) \\
\boldsymbol{\Sigma}_{z|y}^{-1} &= \boldsymbol{\Sigma}_z^{-1} + \mathbf{W}^\top\boldsymbol{\Sigma}_y^{-1}\mathbf{W} \\
\boldsymbol{\mu}_{z|y} &= \boldsymbol{\Sigma}_{z|y} [\mathbf{W}^\top\boldsymbol{\Sigma}_y^{-1}(\mathbf{y} - \mathbf{b}) + \boldsymbol{\Sigma}_z^{-1}\boldsymbol{\mu}_z]
\end{aligned}$$

This is known as **Bayes rule for Gaussians**. Furthermore, the normalization constant of the posterior is given by

$$p(\mathbf{y}) = \int \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_z)\mathcal{N}(\mathbf{y}|\mathbf{W}\mathbf{z} + \mathbf{b}, \boldsymbol{\Sigma}_y) d\mathbf{z} = \mathcal{N}(\mathbf{y}|\mathbf{W}\boldsymbol{\mu}_z + \mathbf{b}, \boldsymbol{\Sigma}_y + \mathbf{W}\boldsymbol{\Sigma}_z\mathbf{W}^\top)$$

The interpretation is as follows: in the generative story ($\mathbf{z} \rightarrow \mathbf{y}$), we first sample $\mathbf{z}$ from the prior $p(\mathbf{z})$, then we sample $\mathbf{y}$ from the likelihood $p(\mathbf{y}|\mathbf{z})$; in the inference story ($\mathbf{y} \rightarrow \mathbf{z}$), we observe $\mathbf{y}$ and want to infer the posterior distribution $p(\mathbf{z}|\mathbf{y})$. While the prior and likelihood are specified by the modeler, the posterior is determined by the data $\mathbf{y}$ through the Bayes rule.

2 Maximum Likelihood Estimation

The process of estimating model parameters θ from data $\mathcal{D}$ is called **model fitting**, or training, and is at the heart of machine learning. The key optimization problem is,

$$\hat{\theta} = \operatorname{argmax}_{\theta} \mathcal{L}(\theta)$$

where $\mathcal{L}(\theta)$ is some kind of loss function or objective function.

The most common approach to parameter estimation is to pick the parameters that assign the highest probability to the training data, called **maximum likelihood estimation** or **MLE**.

$$\hat{\theta}_{mle} = \operatorname{argmax}_{\theta} p(\mathcal{D} \mid \theta)$$

Often, we assume the training examples are independently sampled from the same distribution (**iid** assumption),

$$p(\mathcal{D} \mid \theta) = \prod_{n=1}^N p(\mathbf{y}_n \mid \mathbf{x}_n, \theta)$$

working with log likelihood

$$\hat{\theta}_{mle} = \operatorname{argmax}_{\theta} \log p(\mathcal{D} \mid \theta) = \operatorname{argmax}_{\theta} \sum_{n=1}^N \log p(\mathbf{y}_n \mid \mathbf{x}_n, \theta)$$

In some tasks, the optimization algorithm is designed to minimize the loss function, so equivalently,

$$\hat{\theta}_{mle} = \operatorname{argmin}_{\theta} \text{NLL}(\theta) = \operatorname{argmin}_{\theta} - \sum_{n=1}^N \log p(\mathbf{y}_n \mid \mathbf{x}_n, \theta)$$

In the special case of unsupervised learning, we have $\mathbf{x}_n = \emptyset$, so the MLE changes $p(\mathbf{y}_n \mid \mathbf{x}_n, \theta) = p(\mathbf{y}_n \mid \theta)$

Justification

To justify the use of the MLE is that the resulting predictive distribution $p(\mathbf{y} \mid \hat{\theta}_{mle})$ is as close as possible to the **empirical distribution** of the data. The empirical distribution is defined by

$$p_{\mathcal{D}}(\mathbf{x}, \mathbf{y}) = p_{\mathcal{D}}(\mathbf{y} \mid \mathbf{x}) p_{\mathcal{D}}(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N \delta(\mathbf{x} - \mathbf{x}_n) \delta(\mathbf{y} - \mathbf{y}_n)$$

the empirical distribution is a series of delta functions at the observed training points.

If we define $q(Y \mid \mathbf{x})$ denotes the model's predictive distribution, then the expected KL then becomes

$$\begin{aligned} \mathbb{E}_{p_{\mathcal{D}}(\mathbf{x})} [D_{\text{KL}}(p_{\mathcal{D}}(Y \mid \mathbf{x}) \parallel q(Y \mid \mathbf{x}))] &= \sum_{\mathbf{x}} p_{\mathcal{D}}(\mathbf{x}) \left[\sum_{\mathbf{y}} p_{\mathcal{D}}(\mathbf{y} \mid \mathbf{x}) \log \frac{p_{\mathcal{D}}(\mathbf{y} \mid \mathbf{x})}{q(\mathbf{y} \mid \mathbf{x})} \right] \\ &= \text{const} - \sum_{\mathbf{x}, \mathbf{y}} p_{\mathcal{D}}(\mathbf{x}, \mathbf{y}) \log q(\mathbf{y} \mid \mathbf{x}) \\ &= \text{const} - \frac{1}{N} \sum_{n=1}^N \log p(\mathbf{y}_n \mid \mathbf{x}_n, \theta) \end{aligned}$$

Minimizing the expected KL divergence is equivalent to minimizing the conditional NLL.

3 Matrix Calculus

3.1 Functions that map scalars to scalars

A scalar-argument function $f : \mathbb{R} \rightarrow \mathbb{R}$, its derivative is defined as

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

It can be interpreted as the slope of the tangent line at $f(x)$, and hence

$$f(x+h) \approx f(x) + f'(x)h$$

3.2 Functions that map vectors to scalars

A vector-argument functions, $f : \mathbb{R}^n \rightarrow \mathbb{R}$, by defining the **partial derivative** of f with respect to x_i to be

$$\frac{\partial f}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(\mathbf{x} + h\mathbf{e}_i) - f(\mathbf{x})}{h}$$

where $\mathbf{e}_i$ is the i 'th unit vector.

The **gradient** of a function at a point $\mathbf{x}$ is the vector of its partial derivatives:

$$\mathbf{g} = \frac{\partial f}{\partial \mathbf{x}} = \nabla f = \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{pmatrix}$$

The **directional derivative** measures how much the function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ changes along a direction $\mathbf{v}$ in space,

$$D_{\mathbf{v}}f(\mathbf{x}) = \lim_{h \rightarrow 0} \frac{f(\mathbf{x} + h\mathbf{v}) - f(\mathbf{x})}{h} = \nabla f^\top \mathbf{v}$$

Suppose that some of the arguments to the function depend on each other $f(t, x(t), y(t))$. The **total derivative** of f wrt t is:

$$\frac{df}{dt} = \frac{\partial f}{\partial t} + \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt}$$

then **total differential**,

$$df = \frac{\partial f}{\partial t} dt + \frac{\partial f}{\partial x} dx + \frac{\partial f}{\partial y} dy$$

This measures how much f changes when we change t , both via the direct effect of t on f , but also indirectly, via the effects of t on x and y .

The **Hessian matrix** is the (symmetric) $n \times n$ matrix of second partial derivatives:

$$\mathbf{H}_f = \frac{\partial^2 f}{\partial \mathbf{x}^2} = \nabla^2 f = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}$$

the Hessian is the Jacobian of the gradient.

Consider a differentiable function $f : \mathbb{R}^n \rightarrow \mathbb{R}$. Here are some useful identities:

$$\frac{\partial(\mathbf{a}^\top \mathbf{x})}{\partial \mathbf{x}} = \mathbf{a}$$

$$\frac{\partial(\mathbf{b}^\top \mathbf{A} \mathbf{x})}{\partial \mathbf{x}} = \mathbf{A}^\top \mathbf{b}$$

$$\frac{\partial(\mathbf{x}^\top \mathbf{A} \mathbf{x})}{\partial \mathbf{x}} = (\mathbf{A} + \mathbf{A}^\top) \mathbf{x}$$

3.3 Functions that map vectors to vectors

A function that maps a vector to another vector, $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, can be written as

$$f(\mathbf{x}) = \begin{pmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{pmatrix}$$

where each component $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ is a scalar-valued function. The **Jacobian matrix** of this function is an $m \times n$ matrix of partial derivatives:

$$\mathbf{J}_f(\mathbf{x}) = \frac{\partial \mathbf{f}}{\partial \mathbf{x}^\top} \triangleq \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{pmatrix} = \begin{pmatrix} \nabla f_1(\mathbf{x})^\top \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \end{pmatrix} \in \mathbb{R}^{m \times n}$$

There are two ways to multiply the Jacobian with a vector:

- Jacobian vector product (JVP): $\mathbf{J}_f(\mathbf{x}) \mathbf{v} = \begin{pmatrix} \nabla f_1(\mathbf{x})^\top \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \end{pmatrix} \mathbf{v} = \begin{pmatrix} \nabla f_1(\mathbf{x})^\top \mathbf{v} \\ \vdots \\ \nabla f_m(\mathbf{x})^\top \mathbf{v} \end{pmatrix}.$
- Vector Jacobian product (VJP): $\mathbf{u}^\top \mathbf{J}_f(\mathbf{x}) = \mathbf{u}^\top \left(\frac{\partial \mathbf{f}}{\partial x_1}, \dots, \frac{\partial \mathbf{f}}{\partial x_n} \right) = \left(\mathbf{u} \cdot \frac{\partial \mathbf{f}}{\partial x_1}, \dots, \mathbf{u} \cdot \frac{\partial \mathbf{f}}{\partial x_n} \right).$

Let $h(\mathbf{x}) = g(f(\mathbf{x}))$, by the chain rule of calculus,

$$\mathbf{J}_h(\mathbf{x}) = \mathbf{J}_g(f(\mathbf{x})) \mathbf{J}_f(\mathbf{x})$$

3.4 Functions that map matrices to scalars

A function $f : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}$,

$$\frac{\partial f}{\partial \mathbf{X}} = \begin{pmatrix} \frac{\partial f}{\partial x_{11}} & \cdots & \frac{\partial f}{\partial x_{1n}} \\ \vdots & & \vdots \\ \frac{\partial f}{\partial x_{m1}} & \cdots & \frac{\partial f}{\partial x_{mn}} \end{pmatrix}$$

Here are some useful identities:

$$\frac{\partial}{\partial \mathbf{X}} (\mathbf{a}^\top \mathbf{X} \mathbf{b}) = \mathbf{a} \mathbf{b}^\top, \quad \frac{\partial}{\partial \mathbf{X}} (\mathbf{a}^\top \mathbf{X}^\top \mathbf{b}) = \mathbf{b} \mathbf{a}^\top$$

$$\begin{aligned}
\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{A}\mathbf{X}\mathbf{B}) &= \mathbf{A}^\top \mathbf{B}^\top, & \frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{X}^\top \mathbf{A}) &= \mathbf{A} \\
\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{X}^{-1} \mathbf{A}) &= -\mathbf{X}^{-\top} \mathbf{A}^\top \mathbf{X}^{-\top}, & \frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{X}^\top \mathbf{A} \mathbf{X}) &= (\mathbf{A} + \mathbf{A}^\top) \mathbf{X} \\
\frac{\partial}{\partial \mathbf{X}} \det(\mathbf{A}\mathbf{X}\mathbf{B}) &= \det(\mathbf{A}\mathbf{X}\mathbf{B}) \mathbf{X}^{-\top}, & \frac{\partial}{\partial \mathbf{X}} \log(\det(\mathbf{X})) &= \mathbf{X}^{-\top}
\end{aligned}$$

4 Bound Optimization

If the optimization goal is to maximize a function $\ell(\boldsymbol{\theta})$ wrt its parameters $\boldsymbol{\theta}$. The idea of bound optimization is to construct a surrogate function $Q(\boldsymbol{\theta}; \boldsymbol{\theta}^t)$ which is a tight lowerbound to $\ell(\boldsymbol{\theta})$ such that:

$$Q(\boldsymbol{\theta}; \boldsymbol{\theta}^t) \leq \ell(\boldsymbol{\theta}) \quad \forall \boldsymbol{\theta}, \boldsymbol{\theta}^t \quad \text{and} \quad Q(\boldsymbol{\theta}^t; \boldsymbol{\theta}^t) = \ell(\boldsymbol{\theta}^t)$$

Then, perform the following iterative update:

$$\boldsymbol{\theta}^{t+1} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} Q(\boldsymbol{\theta}; \boldsymbol{\theta}^t)$$

which guarantees monotonic increase of the objective function $\ell(\boldsymbol{\theta})$:

$$\ell(\boldsymbol{\theta}^{t+1}) \geq Q(\boldsymbol{\theta}^{t+1}; \boldsymbol{\theta}^t) \geq Q(\boldsymbol{\theta}^t; \boldsymbol{\theta}^t) = \ell(\boldsymbol{\theta}^t)$$




Figure 1: The solid blue curve is the lower bound, evaluated at $\boldsymbol{\theta}^t$; this touches the objective function at $\boldsymbol{\theta}^t$. We then set $\boldsymbol{\theta}^{t+1}$ to the maximum of the lower bound (blue curve), and fit a new bound at that point (dotted green curve). The maximum of this new bound becomes $\boldsymbol{\theta}^{t+2}$, etc

4.1 EM algorithm

Expectation maximization (EM) algorithm is a special case of bound optimization designed to compute the MLE or MAP parameter estimates for probability models that have **missing data** and/or **hidden variables**.

EM algorithm alternate between:

- E-step: estimating the hidden variables $\mathbf{z}_n$
- M-step: using the fully observed data $\mathbf{y}_n$ to compute the MLE

Lower bound: ELBO

The goal of EM is to maximize the log likelihood of the observed data:

$$\ell(\boldsymbol{\theta}) = \sum_{n=1}^N \log p(\mathbf{y}_n | \boldsymbol{\theta}) = \sum_{n=1}^N \log \left[\sum_{\mathbf{z}_n} p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta}) \right]$$

which is hard to optimize since the log cannot be pushed inside the sum.

Consider a set of arbitrary distribution $q(\mathbf{z}_n)$ over each hidden variable $\mathbf{z}_n$, then apply Jensen's inequality,

$$\begin{aligned}\ell(\boldsymbol{\theta}) &= \sum_{n=1}^N \log \left[\sum_{\mathbf{z}_n} q(\mathbf{z}_n) \frac{p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})}{q(\mathbf{z}_n)} \right] \geq \sum_{n=1}^N \sum_{\mathbf{z}_n} q(\mathbf{z}_n) \log \frac{p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})}{q(\mathbf{z}_n)} \\ &= \sum_n \mathbb{E}_{q_n} [\log p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})] + \mathbb{H}(q_n) \\ &= \sum_n L(\boldsymbol{\theta}, q_n) \triangleq L(\boldsymbol{\theta}, \{q_n\})\end{aligned}$$

where $L(\boldsymbol{\theta}, \{q_n\})$ is the evidence lower bound (ELBO).

E step: fix $\boldsymbol{\theta} = \boldsymbol{\theta}^t$, and update $\{q_n\}$

The lower bound is a sum of N terms, each of which has the form:

$$\begin{aligned}L(\boldsymbol{\theta}, q_n) &= \sum_{\mathbf{z}_n} q_n(\mathbf{z}_n) \log \frac{p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})}{q_n(\mathbf{z}_n)} \\ &= \sum_{\mathbf{z}_n} q_n(\mathbf{z}_n) \log \frac{p(\mathbf{z}_n | \mathbf{y}_n, \boldsymbol{\theta}) p(\mathbf{y}_n | \boldsymbol{\theta})}{q_n(\mathbf{z}_n)} \\ &= \sum_{\mathbf{z}_n} q_n(\mathbf{z}_n) \log \frac{p(\mathbf{z}_n | \mathbf{y}_n, \boldsymbol{\theta})}{q_n(\mathbf{z}_n)} + \sum_{\mathbf{z}_n} q_n(\mathbf{z}_n) \log p(\mathbf{y}_n | \boldsymbol{\theta}) \\ &= -D_{\text{KL}}(q_n(\mathbf{z}_n) \| p(\mathbf{z}_n | \mathbf{y}_n, \boldsymbol{\theta})) + \log p(\mathbf{y}_n | \boldsymbol{\theta})\end{aligned}$$

where $D_{\text{KL}}(q \| p) \geq 0$ and $D_{\text{KL}}(q \| p) = 0$ iff $q = p$. Hence, the maximization of $L(\boldsymbol{\theta}, \{q_n\})$ wrt $\{q_n\}$ can be done by setting each one to $q_n^* = p(\mathbf{z}_n | \mathbf{y}_n, \boldsymbol{\theta})$. Note that, $p(\mathbf{z}_n | \mathbf{y}_n, \boldsymbol{\theta}) = \frac{p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})}{\sum_{\mathbf{z}'_n} p(\mathbf{y}_n, \mathbf{z}'_n | \boldsymbol{\theta})}$ and $p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta}) = p(\mathbf{y}_n | \mathbf{z}_n, \boldsymbol{\theta}) p(\mathbf{z}_n)$ are computable since the prior and conditional distribution on $\mathbf{z}_n$ is defined when defining the probabilistic latent model.

This can be shown that fits into the $Q(\boldsymbol{\theta}, \boldsymbol{\theta}^t)$ formulation.

M step: fix $\{q_n\} = \{q_n^*\}$, and update $\boldsymbol{\theta}$

In the M step, we need to maximize $L(\boldsymbol{\theta}, \{q_n^t\})$ wrt $\boldsymbol{\theta}$, where the q_n^t are the distributions computed in the E step at iteration t . Since the entropy terms $\mathbb{H}(q_n)$ are constant wrt $\boldsymbol{\theta}$, so dropping them in the M step. Then, left with

$$\ell^t(\boldsymbol{\theta}) = \sum_n \mathbb{E}_{q_n^t} [\log p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})]$$

This is the **expected complete data log likelihood**. Then, the optimization task is,

$$\boldsymbol{\theta}^{t+1} = \arg \max_{\boldsymbol{\theta}} \sum_n \mathbb{E}_{q_n^t} [\log p(\mathbf{y}_n, \mathbf{z}_n | \boldsymbol{\theta})]$$

5 Mixture Models

To create more complex probability models is to take a convex combination of simple distributions. This is called a mixture model.

$$p(\mathbf{y}|\boldsymbol{\theta}) = \sum_{k=1}^K p(z = k|\boldsymbol{\theta})p(\mathbf{y}|z = k, \boldsymbol{\theta}) = \sum_{k=1}^K \pi_k p(\mathbf{y}|z = k, \boldsymbol{\theta})$$

π_k is the mixture weight for component k , and $\sum_{k=1}^K \pi_k = 1$, and the **latent variable** $z \in \{1, \dots, K\}$ which specifies which distribution to use for generating the output $\mathbf{y}$.

5.1 Gaussian mixture models

A **Gaussian mixture models (GMM)** is a mixture model where each component is a Gaussian distribution,

$$p(\mathbf{y}|\boldsymbol{\theta}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{y}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

GMMs are often used for unsupervised clustering of real-valued data samples $\mathbf{y}_n \in \mathbb{R}^D$. First, the parameters $\boldsymbol{\theta}$ of the GMM can be estimated using MLE, often with the Expectation-Maximization (EM) algorithm.

$$\hat{\boldsymbol{\theta}} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} \log p(\{\mathbf{y}_n : n = 1, \dots, N\} | \boldsymbol{\theta})$$

Then, we associate each data point $\mathbf{y}_n$ with a discrete latent variable $z_n \in \{1, \dots, K\}$, which indicates which mixture component or cluster generated the data point. The posterior distribution over the latent variable is given by Bayes' rule:

$$p(z_n = k|\mathbf{y}_n, \boldsymbol{\theta}) = \frac{p(\mathbf{y}_n|z_n = k, \boldsymbol{\theta})p(z_n = k|\boldsymbol{\theta})}{\sum_{j=1}^K p(\mathbf{y}_n|z_n = j, \boldsymbol{\theta})p(z_n = j|\boldsymbol{\theta})}$$

This quantity is known as the **responsibility** of cluster k for data point $\mathbf{y}_n$. Then, we can compute the most probable cluster assignment as:

$$\hat{z}_n = \underset{k}{\operatorname{argmax}} p(z_n = k|\mathbf{y}_n, \boldsymbol{\theta})$$

This forms the basis of the clustering algorithms like the K-means clustering.

5.2 EM for a GMM

E step

The E step computes the responsibility of cluster k for generating data point n , as estimated using the current parameter estimates $\boldsymbol{\theta}^t$:

$$r_{nk}^t = p^*(z_n = k|\mathbf{y}_n, \boldsymbol{\theta}^t) = \frac{\pi_k^t p(\mathbf{y}_n|z_n = k, \boldsymbol{\theta}^t)}{\sum_{k'} \pi_{k'}^t p(\mathbf{y}_n|z_n = k', \boldsymbol{\theta}^t)}$$

M step

The M step maximizes the expected complete data log likelihood, given by

$$\begin{aligned}
\ell^t(\boldsymbol{\theta}) &= \mathbb{E} \left[\sum_n \log p(z_n | \boldsymbol{\pi}) + \sum_n \log p(\mathbf{y}_n | z_n, \boldsymbol{\theta}) \right] \\
&= \mathbb{E} \left[\sum_n \log \left(\prod_k \pi_k^{z_{nk}} \right) + \sum_n \log \left(\prod_k \mathcal{N}(\mathbf{y}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)^{z_{nk}} \right) \right] \\
&= \sum_n \sum_k \mathbb{E}[z_{nk}] \log \pi_k + \sum_n \sum_k \mathbb{E}[z_{nk}] \log \mathcal{N}(\mathbf{y}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \\
&= \sum_n \sum_k r_{nk}^{(t)} \log(\pi_k) - \frac{1}{2} \sum_n \sum_k r_{nk}^{(t)} [\log |\boldsymbol{\Sigma}_k| + (\mathbf{y}_n - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1} (\mathbf{y}_n - \boldsymbol{\mu}_k)] + \text{const}
\end{aligned}$$

where $z_{nk} = \mathbb{I}(z_n = k)$ is a one-hot encoding of the categorical value z_n . Then, the parameter estimates are

$$\begin{aligned}
\boldsymbol{\mu}_k^{(t+1)} &= \frac{\sum_n r_{nk}^{(t)} \mathbf{y}_n}{r_k^{(t)}} \\
\boldsymbol{\Sigma}_k^{(t+1)} &= \frac{\sum_n r_{nk}^{(t)} (\mathbf{y}_n - \boldsymbol{\mu}_k^{(t+1)}) (\mathbf{y}_n - \boldsymbol{\mu}_k^{(t+1)})^\top}{r_k^{(t)}} \\
&= \frac{\sum_n r_{nk}^{(t)} \mathbf{y}_n \mathbf{y}_n^\top}{r_k^{(t)}} - \boldsymbol{\mu}_k^{(t+1)} (\boldsymbol{\mu}_k^{(t+1)})^\top
\end{aligned}$$

where $r_k^{(t)} \triangleq \sum_n r_{nk}^{(t)}$ is the weighted number of points assigned to cluster k . The mean of cluster k is just the weighted average of all points assigned to cluster k , and the covariance is proportional to the weighted empirical scatter matrix.

The M step for the mixture weights is a weighted form of the usual MLE:

$$\pi_k^{(t+1)} = \frac{1}{N} \sum_n r_{nk}^{(t)} = \frac{r_k^{(t)}}{N}$$

5.3 Bernoulli mixture model

A **Bernoulli mixture model (BMM)** can be used for binary-valued data, where each mixture component has the following form:

$$p(\mathbf{y} | z = k, \boldsymbol{\theta}) = \prod_{d=1}^D \text{Ber}(y_d | \mu_{dk}) = \prod_{d=1}^D \mu_{dk}^{y_d} (1 - \mu_{dk})^{1-y_d}$$

Part II

Linear Models and Neural Networks

6 Generalized Linear Models

A GLM is a conditional version of an exponential family distribution $p(y | \theta) = h(y) \exp(\theta^\top T(y) - A(\theta))$, in which the natural parameters are a linear function of the input. More precisely, the model has the following form:

$$p(y_n | \mathbf{x}_n, \mathbf{w}, \sigma^2) = \exp \left[\frac{y_n \eta_n - A(\eta_n)}{\sigma^2} + \log h(y_n, \sigma^2) \right]$$

where

- y_n is the output/response variable, $\mathbf{x}_n$ is the input/feature vector, $\mathbf{w}$ and σ is the learnable model parameter
- $\eta_n \triangleq \mathbf{w}^\top \mathbf{x}_n$ is the (input dependent) natural parameter or canonical parameter
- $A(\eta_n)$ is the log normalizer
- $\mathcal{T}(y_n) = y_n$ is the sufficient statistic, which is the smallest summary of the output data y that retains everything relevant for estimating the parameter θ – while it may still lose other information about y itself.
- σ^2 is the dispersion term
- $h(y_n, \sigma^2)$ is the scaling constant or base measure

We can show that the mean and variance of the response variable are as follows:

$$\int p(y_n | \eta_n, \sigma^2) dy_n = 1 \quad \text{for all } \eta_n$$

To derive the mean, differentiate the equation w.r.t to η_n

$$\begin{aligned} \int \frac{\partial p(y_n | \eta_n, \sigma^2)}{\partial \eta_n} dy_n &= 0 \\ \int p(y_n | \eta_n, \sigma^2) \cdot \frac{y_n - A'(\eta_n)}{\sigma^2} dy_n &= 0 \\ \frac{1}{\sigma^2} \left[\int y_n p(y_n | \eta_n, \sigma^2) dy_n - A'(\eta_n) \right] &= 0 \end{aligned}$$

thus, the **mean** is:

$$\mathbb{E}[y_n | \mathbf{x}_n, \mathbf{w}, \sigma^2] = A'(\eta_n) \triangleq \ell^{-1}(\eta_n)$$

where the mapping from the linear inputs to the mean of the output using $\mu_n = \ell^{-1}(\eta_n)$, and the function ℓ is known as the link function, and ℓ^{-1} is known as the mean function.

To derive the variance, differentiate the above equation w.r.t η_n again,

$$\begin{aligned} \int \left[-A''(\eta_n) \cdot p + (y_n - A'(\eta_n)) \cdot p \cdot \frac{y_n - A'(\eta_n)}{\sigma^2} \right] dy_n &= 0 \\ -A''(\eta_n) \int p dy_n + \frac{1}{\sigma^2} \int (y_n - A'(\eta_n))^2 p dy_n &= 0 \end{aligned}$$

thus, the **variance** is:

$$\mathbb{V}[y_n | \mathbf{x}_n, \mathbf{w}, \sigma^2] = A''(\eta_n) \sigma^2$$

6.1 Logistic regression

Binary logistic regression is used when $y_n \in \{0, 1\}$. It has the form

$$p(y_n | \mathbf{x}_n, \mathbf{w}) = \mu_n^{y_n} (1 - \mu_n)^{1-y_n}$$

where

$$\mu_n = \sigma(\eta_n) = \frac{1}{1 + \exp(-\eta_n)}, \quad \eta_n = \mathbf{w}^\top \mathbf{x}_n.$$

Hence

$$\log p(y_n | \mathbf{x}_n, \mathbf{w}) = y_n \log \mu_n + (1 - y_n) \log(1 - \mu_n)$$

Using $\mu_n = \exp(\eta_n)/(1 + \exp(\eta_n))$, in GLM form:

$$\log p(y_n | \mathbf{x}_n, \mathbf{w}) = y_n \eta_n - \log(1 + \exp(\eta_n))$$

We see that $A(\eta_n) = \log(1 + \exp(\eta_n))$ and hence

$$\mathbb{E}[y_n] = A'(\eta_n) = \frac{1}{1 + \exp(-\eta_n)}, \quad \mathbb{V}[y_n] = A''(\eta_n) = \mu_n(1 - \mu_n)$$

Thus the link function is the logit function, $\ell(\mu_n) = \log\left(\frac{\mu_n}{1-\mu_n}\right)$, and the mean function is the sigmoid function, $\ell^{-1}(\eta_n) = \sigma(\eta_n)$.

6.2 Linear regression

Linear regression has the form with $y_n \in \mathbb{R}$

$$p(y_n | \mathbf{x}_n, \mathbf{w}, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(y_n - \mathbf{w}^\top \mathbf{x}_n)^2\right)$$

Hence

$$\log p(y_n | \mathbf{x}_n, \mathbf{w}, \sigma^2) = -\frac{1}{2\sigma^2}(y_n - \eta_n)^2 - \frac{1}{2} \log(2\pi\sigma^2)$$

where $\eta_n = \mathbf{w}^\top \mathbf{x}_n$. In GLM form:

$$\log p(y_n | \mathbf{x}_n, \mathbf{w}, \sigma^2) = \frac{y_n \eta_n - \frac{\eta_n^2}{2}}{\sigma^2} - \frac{1}{2} \left(\frac{y_n^2}{\sigma^2} + \log(2\pi\sigma^2) \right)$$

We see that $A(\eta_n) = \eta_n^2/2$ and hence

$$\mathbb{E}[y_n] = \eta_n = \mathbf{w}^\top \mathbf{x}_n, \quad \mathbb{V}[y_n] = \sigma^2$$

6.3 Binomial regression

If the response variable is the number of successes in N_n trials, $y_n \in \{0, \dots, N_n\}$, we can use binomial regression, which is defined by

$$p(y_n | \mathbf{x}_n, N_n, \mathbf{w}) = \text{Bin}(y_n | \sigma(\mathbf{w}^\top \mathbf{x}_n), N_n)$$

note that binary logistic regression is the special case when $N_n = 1$.

The log pdf is given by

$$\begin{aligned}\log p(y_n | \mathbf{x}_n, N_n, \mathbf{w}) &= y_n \log \mu_n + (N_n - y_n) \log(1 - \mu_n) + \log \binom{N_n}{y_n} \\ &= y_n \log \left(\frac{\mu_n}{1 - \mu_n} \right) + N_n \log(1 - \mu_n) + \log \binom{N_n}{y_n}\end{aligned}$$

where $\mu_n = \sigma(\eta_n)$. To rewrite this in GLM form, let us define

$$\eta_n \triangleq \log \left(\frac{\mu_n}{1 - \mu_n} \right) = \log \left[\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_n}} \frac{1 + e^{-\mathbf{w}^\top \mathbf{x}_n}}{e^{-\mathbf{w}^\top \mathbf{x}_n}} \right] = \log \left(\frac{1}{e^{-\mathbf{w}^\top \mathbf{x}_n}} \right) = \mathbf{w}^\top \mathbf{x}_n$$

Hence we can write binomial regression in GLM form as follows:

$$\log p(y_n | \mathbf{x}_n, N_n, \mathbf{w}) = y_n \eta_n - A(\eta_n) + h(y_n)$$

where

$$h(y_n) = \log \binom{N_n}{y_n}, \quad A(\eta_n) = -N_n \log(1 - \mu_n) = N_n \log(1 + e^{\eta_n})$$

Hence,

$$\mathbb{E}[y_n] = \frac{dA}{d\eta_n} = \frac{N_n e^{\eta_n}}{1 + e^{\eta_n}} = N_n \mu_n, \quad \mathbb{V}[y_n] = \frac{d^2 A}{d\eta_n^2} = N_n \mu_n (1 - \mu_n)$$

6.4 Poisson regression

If the response variable is an integer count, $y_n \in \{0, 1, \dots\}$, we can use Poisson regression, which is defined by

$$p(y_n | \mathbf{x}_n, \mathbf{w}) = \text{Poi}(y_n | \exp(\mathbf{w}^\top \mathbf{x}_n))$$

where

$$\text{Poi}(y | \mu) = e^{-\mu} \frac{\mu^y}{y!}$$

is the Poisson distribution.

The log pdf is given by

$$\log p(y_n | \mathbf{x}_n, \mathbf{w}) = y_n \log \mu_n - \mu_n - \log(y_n!)$$

where $\mu_n = \exp(\mathbf{w}^\top \mathbf{x}_n)$. Hence in GLM form we have

$$\log p(y_n | \mathbf{x}_n, \mathbf{w}) = y_n \eta_n - A(\eta_n) + h(y_n)$$

where $\eta_n = \log(\mu_n) = \mathbf{w}^\top \mathbf{x}_n$, $A(\eta_n) = \mu_n = e^{\eta_n}$, and $h(y_n) = -\log(y_n!)$. Hence

$$\mathbb{E}[y_n] = \frac{dA}{d\eta_n} = e^{\eta_n} = \mu_n, \quad \mathbb{V}[y_n] = \frac{d^2 A}{d\eta_n^2} = e^{\eta_n} = \mu_n$$

6.5 Canonical and non-canonical link functions

Link functions connect three quantities:

- θ_n — the **natural parameter** of the *exponential family distribution*.
- $\mu_n = \mathbb{E}[y_n]$ — the **mean** of the output distribution.
- $\eta_n = \mathbf{w}^\top \mathbf{x}_n$ — the **linear predictor**, built from the *inputs* and weights.

The two connecting functions:

$$\theta_n \xrightarrow{A'} \mu_n \xrightarrow{\ell} \eta_n$$

- A' (derivative of the log-partition function) always maps $\theta_n \rightarrow \mu_n$. This relationship is fixed by the choice of exponential-family distribution (Bernoulli, Gaussian, Poisson, ...) and never changes.
- ℓ , the **link function**, maps $\mu_n \rightarrow \eta_n$. This is the function *you choose* when designing the GLM. Its inverse, ℓ^{-1} , is the **mean function**: $\mu_n = \ell^{-1}(\eta_n)$ — the direction actually used to generate predictions from η_n .

Canonical link. Choosing $\ell = (A')^{-1}$ makes the link exactly undo A' :

$$\eta_n = g(\mu_n) = (A')^{-1}(A'(\theta_n)) = \theta_n$$

so the natural parameter and linear predictor coincide: $\theta_n = \eta_n$.

Non-canonical link. Choosing any $g \neq (A')^{-1}$ gives

$$\eta_n = g(A'(\theta_n)) \neq \theta_n \quad \text{in general.}$$

The mean is still recovered correctly via $\mu_n = g^{-1}(\eta_n)$, but the natural parameter θ_n is no longer equal to the linear predictor η_n —recovering it requires the separate map $(A')^{-1}(\mu_n)$.

Examples of links:

- Logit: canonical link $\theta_n = \ell(\mu_n) = \log \frac{\mu_n}{1-\mu_n}$.
- Complementary log-log: non-canonical link $\eta_n = \ell(\mu_n) = \log(-\log(1 - \mu_n))$

6.6 Maximum likelihood estimation

GLMs can be fit the negative log-likelihood as the objective function (ignoring constant term h):

$$\text{NLL}(\mathbf{w}) = -\log p(\mathcal{D} \mid \mathbf{w}) = -\frac{1}{\sigma^2} \sum_{n=1}^N \ell_n = -\frac{1}{\sigma^2} \sum_{n=1}^N [\eta_n y_n - A(\eta_n)]$$

where $\eta_n = \mathbf{w}^\top \mathbf{x}_n$.

To find the MLE, we must solve $\nabla \text{NLL}(\mathbf{w}) = \mathbf{g}(\mathbf{w}) = 0$. For notational simplicity, we will assume $\sigma^2 = 1$. The gradient for a single term is:

$$\begin{aligned} \mathbf{g}_n &= -\frac{1}{\sigma^2} \frac{\partial \ell_n}{\partial \mathbf{w}} = -\frac{1}{\sigma^2} \frac{\partial \ell_n}{\partial \eta_n} \frac{\partial \eta_n}{\partial \mathbf{w}} \\ &= -\frac{1}{\sigma^2} (y_n - A'(\eta_n)) \mathbf{x}_n = -\frac{1}{\sigma^2} (y_n - \mu_n) \mathbf{x}_n \end{aligned}$$

where $\mu_n = f(\mathbf{w}^\top \mathbf{x})$, and f is the inverse link function that maps from canonical parameters to mean parameters. If we interpret $(y_n - \mu_n)$ as an error signal, then the gradient $-\frac{1}{\sigma^2} \sum_{n=1}^N (y_n - \mu_n) \mathbf{x}_n$ weights each input $\mathbf{x}_n$ by its error, and averages the result.

Gradient-based optimization will find a stationary points where $\mathbf{g}(\mathbf{w}) = 0$, we need to examine the Hessian matrix to determine whether it is the global optimum. The Hessian is given by

$$\begin{aligned} \mathbf{H} &= \frac{\partial^2}{\partial \mathbf{w} \partial \mathbf{w}^\top} \text{NLL}(\mathbf{w}) \\ &= - \sum_{n=1}^N \frac{\partial \mathbf{g}_n}{\partial \mathbf{w}^\top} = - \sum_{n=1}^N \frac{\partial \mathbf{g}_n}{\partial \mu_n} \frac{\partial \mu_n}{\partial \mathbf{w}^\top} = - \sum_{n=1}^N \mathbf{x}_n f'(\mathbf{w}^\top \mathbf{x}_n) \mathbf{x}_n^\top \end{aligned}$$

In general, the Hessian is positive definite, since $f'(\mathbf{w}^\top \mathbf{x}_n) > 0$, hence the negative log-likelihood is convex, so the MLE for a GLM is unique given $f(\mathbf{w}^\top \mathbf{x}_n) > 0$ for all n .

Based on these results, gradient-based optimizers can be used to estimate the model weights.

Stochastic gradient descent

The goal is to solve

$$\hat{\mathbf{w}} \triangleq \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \text{NLL}(\mathbf{w})$$

One of the methods is stochastic gradient descent: we use a minibatch from the whole dataset of size N , then we get the following update equation:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha_t \nabla_{\mathbf{w}} \text{NLL}(\mathbf{w}_t) = \mathbf{w}_t - \alpha_t (\mu_n - y_n) \mathbf{x}_n$$

where we replaced the average over all N samples with a stochastically chosen minibatch size at each step t .

Since we know the objective is convex, then it can be shown that this will converge to a global optimum given the learning rate α_t is decayed appropriately at each step t .

Part III

Dimentionality Reduction

7 Principal Component Analysis

PCA algorithm is to find a **linear and orthogonal** projection of the high dimensional data $x \in \mathbb{R}^D$, to a low dimensional subspace $z \in \mathbb{R}^L$, such that the low dimensional representation is a "good approximation" to the original data. If we project or encode $\mathbf{x}$ to get $\mathbf{z} = \mathbf{W}^T \mathbf{x}$, and then unproject or decode $\mathbf{z}$ to get $\hat{\mathbf{x}} = \mathbf{W}\mathbf{z}$, then we want $\hat{\mathbf{x}}$ to be close to $\mathbf{x}$ in ℓ_2 distance. In particular, we can define the following **reconstruction error**:

$$\mathcal{L}(\mathbf{W}) \triangleq \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \text{decode}(\text{encode}(\mathbf{x}_n; \mathbf{W}); \mathbf{W})\|_2^2$$

PCA algorithm states that this objective can be minimized by setting $\mathbf{W} = \mathbf{U}_L$, where $\mathbf{U}_L$ contains the the L eigenvectors with lages eigenvalues of the empirical covariance matrix

$$\hat{\Sigma} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T = \frac{1}{N} \mathbf{X}_c^T \mathbf{X}_c$$

where $\mathbf{X}_c$ is a centered version of the $N \times D$ design matrix.

7.1 PCA setup

Suppose we have an (unlabeled) dataset $\mathcal{D} = \{\mathbf{x}_n : n = 1 : N\}$, where $\mathbf{x}_n \in \mathbb{R}^D$. We can represent this as an $N \times D$ data matrix $\mathbf{X}$. We will assume $\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \mathbf{0}$, which we can ensure by centering the data.

We would like to approximate each $\mathbf{x}_n$ by a low dimensional representation, $\mathbf{z}_n \in \mathbb{R}^L$. We assume that each $\mathbf{x}_n$ can be explained in terms of a weighted combination of basis functions $\mathbf{w}_1, \dots, \mathbf{w}_L$, where each $\mathbf{w}_k \in \mathbb{R}^D$, and where the weights are given by $\mathbf{z}_n \in \mathbb{R}^L$, i.e.

$$\mathbf{x}_n \approx \sum_{k=1}^L z_{nk} \mathbf{w}_k$$

The vector $\mathbf{z}_n$ is the low dimensional representation of $\mathbf{x}_n$, and is known as the **latent vector**, the collection of these latent variables are called the **latent factors**.

The error produced by this approximation:

$$\mathcal{L}(\mathbf{W}, \mathbf{Z}) = \frac{1}{N} \|\mathbf{X} - \mathbf{Z}\mathbf{W}^T\|_F^2 = \frac{1}{N} \|\mathbf{X}^T - \mathbf{W}\mathbf{Z}^T\|_F^2 = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{W}\mathbf{z}_n\|^2$$

where

- $\mathbf{X} \in \mathbb{R}^{N \times D}$: data matrix, rows are centered data points $\mathbf{x}_n^T$
- $\mathbf{Z} \in \mathbb{R}^{N \times L}$: latent code matrix, rows are centered data points $\mathbf{z}_n^T$
- $\mathbf{W} \in \mathbb{R}^{D \times L}$: basis matrix, columns $\mathbf{w}_1, \dots, \mathbf{w}_L$ are orthonormal
- approximating each $\mathbf{x}_n$ by $\hat{\mathbf{x}}_n = \mathbf{W}\mathbf{z}_n$; or, $\hat{\mathbf{X}} = \mathbf{Z}\mathbf{W}^T$.
- the minimization of $\mathcal{L}$ is subject to constraints $\mathbf{w}_k^T \mathbf{w}_j = 0$ for $k \neq j$; 1 for $k = j$, i.e., $\mathbf{W}$ is orthogonal

7.2 Algorithm derivation

Base case

Starting by considering the best 1d solution, $\mathbf{w}_1 \in \mathbb{R}^D$, letting the coefficients for each data-points associated with the first basis vector be denoted by $\tilde{\mathbf{z}}_1 = [z_{11}, \dots, z_{N1}] \in \mathbb{R}^N$. The reconstruction error is

$$\begin{aligned}\mathcal{L}(\mathbf{w}_1, \tilde{\mathbf{z}}_1) &= \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - z_{n1} \mathbf{w}_1\|^2 = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - z_{n1} \mathbf{w}_1)^\top (\mathbf{x}_n - z_{n1} \mathbf{w}_1) \\ &= \frac{1}{N} \sum_{n=1}^N [\mathbf{x}_n^\top \mathbf{x}_n - 2z_{n1} \mathbf{w}_1^\top \mathbf{x}_n + z_{n1}^2 \mathbf{w}_1^\top \mathbf{w}_1] \\ &= \frac{1}{N} \sum_{n=1}^N [\mathbf{x}_n^\top \mathbf{x}_n - 2z_{n1} \mathbf{w}_1^\top \mathbf{x}_n + z_{n1}^2] \quad \text{since } \mathbf{w}_1^\top \mathbf{w}_1 = 1\end{aligned}$$

Taking derivatives wrt z_{n1} and equating to zero gives

$$\frac{\partial}{\partial z_{n1}} \mathcal{L}(\mathbf{w}_1, \tilde{\mathbf{z}}_1) = \frac{1}{N} [-2\mathbf{w}_1^\top \mathbf{x}_n + 2z_{n1}] = 0 \Rightarrow z_{n1} = \mathbf{w}_1^\top \mathbf{x}_n$$

So the optimal embedding is obtained by orthogonally projecting the data onto $\mathbf{w}_1$. Plugging this back in gives the loss for the weights:

$$\mathcal{L}(\mathbf{w}_1) = \mathcal{L}(\mathbf{w}_1, \tilde{\mathbf{z}}_1^*(\mathbf{w}_1)) = \frac{1}{N} \sum_{n=1}^N [\mathbf{x}_n^\top \mathbf{x}_n - z_{n1}^2] = \text{const} - \frac{1}{N} \sum_{n=1}^N z_{n1}^2$$

To solve for $\mathbf{w}_1$, note that

$$\mathcal{L}(\mathbf{w}_1) = -\frac{1}{N} \sum_{n=1}^N z_{n1}^2 = -\frac{1}{N} \sum_{n=1}^N \mathbf{w}_1^\top \mathbf{x}_n \mathbf{x}_n^\top \mathbf{w}_1 = -\mathbf{w}_1^\top \hat{\Sigma} \mathbf{w}_1$$

where Σ is the empirical covariance matrix (since we assumed the data is centered). We can trivially optimize this by letting $\|\mathbf{w}_1\| \rightarrow \infty$, so we impose the constraint $\|\mathbf{w}_1\| = 1$ and instead optimize

$$\tilde{\mathcal{L}}(\mathbf{w}_1) = \mathbf{w}_1^\top \hat{\Sigma} \mathbf{w}_1 - \lambda_1 (\mathbf{w}_1^\top \mathbf{w}_1 - 1)$$

where λ_1 is a Lagrange multiplier. Taking derivatives and equating to zero we have

$$\frac{\partial}{\partial \mathbf{w}_1} \tilde{\mathcal{L}}(\mathbf{w}_1) = 2\hat{\Sigma} \mathbf{w}_1 - 2\lambda_1 \mathbf{w}_1 = 0$$

$$\hat{\Sigma} \mathbf{w}_1 = \lambda_1 \mathbf{w}_1$$

Hence the optimal direction onto which we should project the data is an eigenvector of the covariance matrix. Left multiplying by $\mathbf{w}_1^\top$ (and using $\mathbf{w}_1^\top \mathbf{w}_1 = 1$) we find

$$\mathbf{w}_1^\top \hat{\Sigma} \mathbf{w}_1 = \lambda_1$$

Since we want to maximize this quantity (minimize the loss), we pick the eigenvector which corresponds to the largest eigenvalue.

Induction Step

For another direction $\mathbf{w}_2$, the reconstruction error is

$$\mathcal{L}(\mathbf{w}_1, \tilde{\mathbf{z}}_1, \mathbf{w}_2, \tilde{\mathbf{z}}_2) = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - z_{n1}\mathbf{w}_1 - z_{n2}\mathbf{w}_2\|^2$$

Optimizing wrt $\mathbf{w}_1$ and $\mathbf{z}_1$ gives the same solution as before. $\frac{\partial \mathcal{L}}{\partial \mathbf{z}_2} = 0$ yields $z_{n2} = \mathbf{w}_2^\top \mathbf{x}_n$. Substituting in yields

$$\mathcal{L}(\mathbf{w}_2) = \frac{1}{N} \sum_{n=1}^N [\mathbf{x}_n^\top \mathbf{x}_n - \mathbf{w}_1^\top \mathbf{x}_n \mathbf{x}_n^\top \mathbf{w}_1 - \mathbf{w}_2^\top \mathbf{x}_n \mathbf{x}_n^\top \mathbf{w}_2] = \text{const} - \mathbf{w}_2^\top \hat{\Sigma} \mathbf{w}_2$$

Dropping the constant term, plugging in the optimal $\mathbf{w}_1$ and adding the constraints yields

$$\tilde{\mathcal{L}}(\mathbf{w}_2) = -\mathbf{w}_2^\top \hat{\Sigma} \mathbf{w}_2 + \lambda_2(\mathbf{w}_2^\top \mathbf{w}_2 - 1) + \lambda_{12}(\mathbf{w}_2^\top \mathbf{w}_1 - 0)$$

with the extra constraint term, it can be shown that the solution is given by the eigenvector with the second largest eigenvalue:

$$\hat{\Sigma} \mathbf{w}_2 = \lambda_2 \mathbf{w}_2$$

The proof continues in this way to show that $\hat{\mathbf{W}} = \mathbf{U}_L$.

7.3 Computational tricks

1. **High-dimensional data.** PCA finds the eigenvectors of the $D \times D$ covariance matrix $\mathbf{X}^\top \mathbf{X}$. If $D > N$, it is faster to work with the $N \times N$ matrix $\mathbf{X}\mathbf{X}^\top$. Let $\mathbf{U}$ be an orthonormal matrix containing the eigenvectors of $\mathbf{X}\mathbf{X}^\top$ with corresponding eigenvalues in Λ . By definition we have $(\mathbf{X}\mathbf{X}^\top)\mathbf{U} = \mathbf{U}\Lambda$. Pre-multiplying by $\mathbf{X}^\top$ gives

$$(\mathbf{X}^\top \mathbf{X})(\mathbf{X}^\top \mathbf{U}) = (\mathbf{X}^\top \mathbf{U})\Lambda$$

from which we see that the eigenvectors of $\mathbf{X}^\top \mathbf{X}$ are $\mathbf{V} = \mathbf{X}^\top \mathbf{U}$, with eigenvalues Λ as before. However, these eigenvectors are not normalized, since $\|\mathbf{v}_j\|^2 = \mathbf{u}_j^\top \mathbf{X}\mathbf{X}^\top \mathbf{u}_j = \lambda_j \mathbf{u}_j^\top \mathbf{u}_j = \lambda_j$. The normalized eigenvectors are given by

$$\mathbf{V} = \mathbf{X}^\top \mathbf{U} \Lambda^{-\frac{1}{2}}$$

2. **Computing PCA using SVD.** Let $\mathbf{U}_\Sigma \Lambda_\Sigma \mathbf{U}_\Sigma^\top$ be the top L eigendecomposition of the covariance matrix $\Sigma \propto \mathbf{X}^\top \mathbf{X}$ (assume $\mathbf{X}$ is centered). Recall the optimal estimate of the projection weights $\mathbf{W}$ is given by the top L eigenvalues, so $\mathbf{W} = \mathbf{U}_\Sigma$. Now let $\mathbf{U}_X \mathbf{S}_X \mathbf{V}_X^\top \approx \mathbf{X}$ be the L -truncated SVD approximation to the data matrix $\mathbf{X}$. By SVD definition, the columns of $\mathbf{V}_X$ are the eigenvectors of $\mathbf{X}^\top \mathbf{X}$, so

$$\mathbf{V}_X = \mathbf{U}_\Sigma = \mathbf{W}.$$

In addition, the eigenvalues of the covariance matrix are related to the singular values of the data matrix via $\lambda_k = s_k^2/N$. Now suppose we are interested in the principal components (or PC scores), rather than the projection matrix. We have

$$\mathbf{Z} = \mathbf{X}\mathbf{W} = \mathbf{U}_X \mathbf{S}_X \mathbf{V}_X^\top \mathbf{V}_X = \mathbf{U}_X \mathbf{S}_X$$

Finally, if we want to approximately reconstruct the data, we have

$$\hat{\mathbf{X}} = \mathbf{Z}\mathbf{W}^\top = \mathbf{U}_X \mathbf{S}_X \mathbf{V}_X^\top$$

This is precisely the same as a **truncated SVD approximation**.

3. **Latent dimentions selection.** Recall that the reconstruction error can be written in terms of eigenvalues

$$\mathcal{L}_L = \frac{1}{|\mathcal{D}|} \sum_{n \in \mathcal{D}} \|\mathbf{x}_n - \hat{\mathbf{x}}_n\|^2 = \sum_{j=L+1}^D \lambda_j$$

where $\hat{\mathbf{x}}_n = \mathbf{W}\mathbf{z}_n + \mu$ and μ is the empirical mean. Thus, as the number of dimentions incrcases, the eigenvalues get smaller, and so does the reconstruction error. Therefore, if we give it more latent dimenions, it will be able to approximate the test date more accurately.

Thus, PCA can be performed either using an eigendecomposition of $\mathbf{\Sigma}$ or an SVD decomposition of $\mathbf{X}$. The latter is often preferable, for computational reasons.

8 Factor Analysis

Factor analysis is a generalization of PCA corresponds to the following linear-Gaussian latent variable probabilistic **generative** model:

$$\mathbf{x} = \mathbf{W}\mathbf{z} + \boldsymbol{\mu} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Psi})$$

$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z} \mid \boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$$

$$p(\mathbf{x} \mid \mathbf{z}, \boldsymbol{\theta}) = \mathcal{N}(\mathbf{x} \mid \mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \boldsymbol{\Psi})$$

where $\mathbf{W}$ is a $D \times L$ matrix, known as the **factor loading matrix**, $\boldsymbol{\Psi}$ is a diagonal $D \times D$ noise covariance matrix, and $\boldsymbol{\mu}$ is a $D \times 1$ mean of observed data

FA can be thought of as a low-rank version of a Gaussian distribution. To see this, note that the induced marginal distribution $p(\mathbf{x} \mid \boldsymbol{\theta})$ is a Gaussian:

$$\begin{aligned} p(\mathbf{x} \mid \boldsymbol{\theta}) &= \int \mathcal{N}(\mathbf{x} \mid \mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \boldsymbol{\Psi}) \mathcal{N}(\mathbf{z} \mid \boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0) d\mathbf{z} \\ &= \mathcal{N}(\mathbf{x} \mid \mathbf{W}\boldsymbol{\mu}_0 + \boldsymbol{\mu}, \boldsymbol{\Psi} + \mathbf{W}\boldsymbol{\Sigma}_0\mathbf{W}^\top) \end{aligned}$$

Hence $\mathbb{E}[\mathbf{x}] = \mathbf{W}\boldsymbol{\mu}_0 + \boldsymbol{\mu}$ and $\text{Cov}[\mathbf{x}] = \mathbf{W}\text{Cov}[\mathbf{z}]\mathbf{W}^\top + \boldsymbol{\Psi} = \mathbf{W}\boldsymbol{\Sigma}_0\mathbf{W}^\top + \boldsymbol{\Psi}$. Without loss of generality, we can set $\boldsymbol{\mu}_0 = \mathbf{0}$ and $\boldsymbol{\Sigma}_0 = \mathbf{I}$. Thus,

$$p(\mathbf{z}) = \mathcal{N}(\mathbf{z} \mid \mathbf{0}, \mathbf{I})$$

$$p(\mathbf{x} \mid \mathbf{z}) = \mathcal{N}(\mathbf{x} \mid \mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \boldsymbol{\Psi})$$

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}, \mathbf{W}\mathbf{W}^\top + \boldsymbol{\Psi})$$

where FA approximates the covariance matrix of the visible vector using a low-rank decomposition $\mathbf{C} = \text{Cov}[\mathbf{x}] = \mathbf{W}\mathbf{W}^\top + \boldsymbol{\Psi}$. The marginal variance of each visible variable is given by $\mathbb{V}[x_d] = \sum_{k=1}^L w_{dk}^2 + \psi_d$ where the first term is the variance due to the common factors, and the second ψ_d term is dimension-specific uniqueness.

Using EM algorithm to estimate the parameters $\mathbf{W}, \boldsymbol{\mu}, \boldsymbol{\Psi}$, then we can compute probabilistic latent embeddings using $p(\mathbf{z} \mid \mathbf{x})$ using Bayes rule for Gaussians,

$$p(\mathbf{z} \mid \mathbf{x}) = \mathcal{N}(\mathbf{z} \mid \mathbf{W}^\top \mathbf{C}^{-1}(\mathbf{x} - \boldsymbol{\mu}), \mathbf{I} - \mathbf{W}^\top \mathbf{C}^{-1} \mathbf{W})$$

8.1 Prerequisite: evidence lower bound (ELBO)

Let data x been generated by some latent variable z through a process $p(z) \rightarrow p(x \mid z)$.

- To compute the marginal likelihood:

$$p(x) = \int p(x \mid z) p(z) dz$$

which tells how well the model explains the data, and should be maximized during training. However, the integral is usually intractable (eg. $p(x \mid z)$ is a neural network).

- To compute posterior $p(z \mid x)$ (i.e. given this data point, what latent value likely produced it). By Bayes' rule,

$$p(z \mid x) = \frac{p(x \mid z)p(z)}{p(x)}$$

and this also requires the intractable $p(x)$ in the denominator.

Since $p(z | x)$ cannot be computed exactly, an approximate posterior $q(z | x)$ is chosen and try to make it as close as possible to the true posterior, measured by KL divergence:

$$D_{KL}(q(z | x) \| p(z | x)) \geq 0$$

expand from the definition of KL divergence and re-arrange,

$$\begin{aligned} \log p(x) &= \underbrace{\mathbb{E}_{q(z|x)}[\log p(x | z)] - D_{KL}(q(z | x) \| p(z))}_{\text{ELBO}} + D_{KL}(q(z | x) \| p(z | x)) \\ &= \mathcal{L}(x) + D_{KL}(q(z|x) \| p(z|x)) \end{aligned}$$

for which,

$$\log p(x) \geq \mathcal{L}(x) \quad \text{since } D_{KL}(\cdot \| \cdot) \geq 0$$

ELBO is an useful variational trick,

- $\log p(x)$ itself is intractable, but $\mathcal{L}(x)$ only involves $q(z | x)$, $p(z)$, and $p(x | z)$ — all things you choose/model directly.
- Maximizing the ELBO does two things simultaneously: it pushes $\log p(x)$ up (better data likelihood), and it pushes $q(z | x)$ closer to the true posterior $p(z | x)$ (since the gap between ELBO and $\log p(x)$ is that KL term).

In the EM algorithm for FA below, expressing ELBO equivalently as $\mathcal{L} = \mathbb{E}_q[\log p(x, z)] - \mathbb{E}_q[\log q(z)]$, according to the lower bound theorem above, maximizing $\log p(x)$ can be done by maximizing $\mathbb{E}_q[\log p(x, z)]$.

8.2 EM algorithm for FA

The FA generative model:

$$p(\mathbf{z}_i) = \mathcal{N}(\mathbf{z}_i | \mathbf{0}, \mathbf{I}), \quad p(\mathbf{x}_i | \mathbf{z}_i) = \mathcal{N}(\mathbf{x}_i | \mathbf{W}\mathbf{z}_i + \boldsymbol{\mu}, \boldsymbol{\Psi})$$

with parameters $\boldsymbol{\theta} = (\mathbf{W}, \boldsymbol{\mu}, \boldsymbol{\Psi})$ shared across all datapoints $i = 1, \dots, N$.

E-step: compute the posterior $p(\mathbf{z}_i | \mathbf{x}_i, \boldsymbol{\theta})$. Because both distributions above are Gaussian and linearly related, this posterior is available in closed form by completing the square:

$$p(\mathbf{z}_i | \mathbf{x}_i, \boldsymbol{\theta}) = \mathcal{N}(\mathbf{z}_i | \mathbf{m}_i, \boldsymbol{\Sigma}_i)$$

$$\boldsymbol{\Sigma}_i = (\mathbf{I}_L + \mathbf{W}^T \boldsymbol{\Psi}^{-1} \mathbf{W})^{-1} \in \mathbb{R}^{L \times L}, \quad \mathbf{m}_i = \boldsymbol{\Sigma}_i \mathbf{W}^T \boldsymbol{\Psi}^{-1} (\mathbf{x}_i - \boldsymbol{\mu}) \in \mathbb{R}^{L \times 1}$$

Note $\boldsymbol{\Sigma}_i$ does not depend on i in practice (no $\mathbf{x}_i$ term appears), so it can be computed once per iteration.

M-step: maximize the expected complete-data log-likelihood $\mathbb{E}[\log p(\mathbf{x}_i, \mathbf{z}_i | \boldsymbol{\theta})]$. It is convenient to estimate $\mathbf{W}$ and $\boldsymbol{\mu}$ jointly by defining

$$\tilde{\mathbf{W}} = (\mathbf{W}, \boldsymbol{\mu}), \quad \tilde{\mathbf{z}} = (\mathbf{z}, 1), \quad \text{so that } \mathbf{x} = \mathbf{W}\mathbf{z} + \boldsymbol{\mu} = \tilde{\mathbf{W}}\tilde{\mathbf{z}}$$

Define:

$$\begin{aligned} \mathbf{b}_i &\triangleq \mathbb{E}[\tilde{\mathbf{z}}_i | \mathbf{x}_i] = [\mathbf{m}_i; 1] \\ \mathbf{C}_i &\triangleq \mathbb{E}[\tilde{\mathbf{z}}_i \tilde{\mathbf{z}}_i^T | \mathbf{x}_i] = \begin{pmatrix} \mathbb{E}[\mathbf{z}_i \mathbf{z}_i^T | \mathbf{x}_i] & \mathbb{E}[\mathbf{z}_i | \mathbf{x}_i] \\ \mathbb{E}[\mathbf{z}_i | \mathbf{x}_i]^T & 1 \end{pmatrix} = \begin{pmatrix} \boldsymbol{\Sigma}_i + \mathbf{m}_i \mathbf{m}_i^T & \mathbf{m}_i \\ \mathbf{m}_i^T & 1 \end{pmatrix} \end{aligned}$$

Maximizing the expected complete-data log-likelihood with respect to $\tilde{\mathbf{W}}$ and Ψ (a weighted-least-squares-type problem, since it is quadratic in $\tilde{\mathbf{W}}$ given $\mathbf{b}_i, \mathbf{C}_i$) yields:

$$\begin{aligned} \arg \max_{\boldsymbol{\theta}} \sum_i \mathbb{E}[\log p(\mathbf{x}_i, \mathbf{z}_i \mid \boldsymbol{\theta})] &= \arg \max_{\boldsymbol{\theta}} \sum_i \mathbb{E}[\log p(\mathbf{z}_i) + \log p(\mathbf{x}_i \mid \mathbf{z}_i, \boldsymbol{\theta})] \\ &= \arg \max_{\boldsymbol{\theta}} \sum_i \mathbb{E}[\log p(\mathbf{x}_i \mid \mathbf{z}_i, \boldsymbol{\theta})] \\ &= \arg \max_{\boldsymbol{\theta}} \left\{ -\frac{N}{2} \log |\Psi| - \frac{1}{2} \sum_{i=1}^N \left[\mathbf{x}_i^\top \Psi^{-1} \mathbf{x}_i - 2 \mathbf{x}_i^\top \Psi^{-1} \tilde{\mathbf{W}} \mathbf{b}_i + \text{tr}(\tilde{\mathbf{W}}^\top \Psi^{-1} \tilde{\mathbf{W}} \mathbf{C}_i) \right] \right\} \end{aligned}$$

Solve for $\boldsymbol{\theta} = (\tilde{\mathbf{W}}, \Psi)$ by setting $\partial Q / \partial \tilde{\mathbf{W}} = 0$ to get $\hat{\tilde{\mathbf{W}}}$, then substituting back and setting $\partial Q / \partial \Psi^{-1} = 0$ to get $\hat{\Psi}$,

$$\begin{aligned} \hat{\tilde{\mathbf{W}}} &= \left[\sum_i \mathbf{x}_i \mathbf{b}_i^\top \right] \left[\sum_i \mathbf{C}_i \right]^{-1} \\ \hat{\Psi} &= \frac{1}{N} \text{diag} \left\{ \sum_i \left(\mathbf{x}_i - \hat{\tilde{\mathbf{W}}} \mathbf{b}_i \right) \mathbf{x}_i^\top \right\} \end{aligned}$$

the $\text{diag}\{\cdot\}$ enforces the modeling assumption that Ψ remains diagonal.

The loop. Alternate E-step (recompute $\mathbf{m}_i, \Sigma_i, \mathbf{b}_i, \mathbf{C}_i$ using current $\boldsymbol{\theta}$) and M-step (recompute $\mathbf{W}, \boldsymbol{\mu}, \Psi$ using current $\mathbf{b}_i, \mathbf{C}_i$) until the marginal log-likelihood converges.

8.3 EM algorithm for PPCA

PPCA is the special case of FA with $\Psi = \sigma^2 \mathbf{I}$. Substituting into the general FA-EM equations:

E-step.

$$\Sigma = (\mathbf{I}_L + \sigma^{-2} \mathbf{W}^T \mathbf{W})^{-1} = \sigma^2 \mathbf{M}^{-1}, \quad \mathbf{M} := \mathbf{W}^T \mathbf{W} + \sigma^2 \mathbf{I}$$

$$\mathbf{m}_i = \mathbf{M}^{-1} \mathbf{W}^T (\mathbf{x}_i - \boldsymbol{\mu})$$

(Σ has no subscript i : it does not depend on $\mathbf{x}_i$, only on $\mathbf{W}, \sigma^2$.)

M-step. The $\tilde{\mathbf{W}}$ update is unchanged in form:

$$\hat{\tilde{\mathbf{W}}} = \left[\sum_i \mathbf{x}_i \mathbf{b}_i^\top \right] \left[\sum_i \mathbf{C}_i \right]^{-1}$$

Since noise is isotropic, σ^2 is a single scalar (average over all D dimensions, rather than a diagonal Ψ):

$$\hat{\sigma}^2 = \frac{1}{ND} \sum_{i=1}^N \text{tr} \left[(\mathbf{x}_i - \hat{\tilde{\mathbf{W}}} \mathbf{b}_i) \mathbf{x}_i^\top \right]$$

8.4 Limiting case: EM algorithm for PCA

Noise-free limit ($\sigma^2 \rightarrow 0$): *EM for PCA.* The posterior collapses to a point estimate ($\Sigma \rightarrow \mathbf{0}$), reducing to alternating least squares. With data centered and column-stacked as $\tilde{\mathbf{X}} \in \mathbb{R}^{D \times N}$:

1. E-step: given current $\mathbf{W}$, project data onto it via least squares $\rightarrow \tilde{\mathbf{Z}}$

$$\tilde{\mathbf{Z}} = (\mathbf{W}^T \mathbf{W})^{-1} \mathbf{W}^T \tilde{\mathbf{X}}$$

2. M-step: given current $\tilde{\mathbf{Z}}$, re-fit $\mathbf{W}$ via least-squares regression of $\mathbf{X}$ onto $\mathbf{Z}$

$$\mathbf{W} = \tilde{\mathbf{X}}\tilde{\mathbf{Z}}^T(\tilde{\mathbf{Z}}\tilde{\mathbf{Z}}^T)^{-1}$$

Why use EM instead of the closed-form eigendecomposition? Each EM iteration only requires $L \times L$ inversions ($L \ll D$ typically), costing $O(NDL + L^3)$ per iteration, versus $O(D^3)$ for eigendecomposing the full covariance matrix $\mathbf{S}$. This matters when D is large (e.g. many neurons) but L is small. EM also extends naturally to missing data, unlike the closed-form solution.

Sanity check. As $\sigma^2 \rightarrow 0$, both EM-derived quantities converge to the standard PCA solution: $\hat{\mathbf{W}} \rightarrow \mathbf{U}_L \mathbf{L}_L^{1/2}$ and $\mathbb{E}[\mathbf{z} \mid \mathbf{x}] \rightarrow (\mathbf{W}^T \mathbf{W})^{-1} \mathbf{W}^T (\mathbf{x} - \bar{\mathbf{x}})$, matching the closed-form Tipping & Bishop result.

8.5 Special case: Probabilistic PCA

Probabilistic PCA (PPCA) is a special case of the factor analysis model in which $\mathbf{W}$ has orthonormal columns and $\Psi = \sigma^2 \mathbf{I}$. Then, the full-form marginal distribution on the visible variables reduces to:

$$\begin{aligned} p(\mathbf{x} \mid \boldsymbol{\mu}, \mathbf{W}, \sigma^2) &= \int \mathcal{N}(\mathbf{x} \mid \mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \Psi) \mathcal{N}(\mathbf{z} \mid \boldsymbol{\mu}_0, \Sigma_0) d\mathbf{z} \\ &= \int \mathcal{N}(\mathbf{x} \mid \mathbf{W}\mathbf{z}, \sigma^2 \mathbf{I}) \mathcal{N}(\mathbf{z} \mid \mathbf{0}, \mathbf{I}) d\mathbf{z} \\ &= \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}, \mathbf{W}\mathbf{W}^T + \sigma^2 \mathbf{I}) = \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}, \mathbf{C}) \end{aligned}$$

Then, the log-likelihood for PPCA is:

$$\begin{aligned} \log p(\mathbf{X} \mid \boldsymbol{\mu}, \mathbf{W}, \sigma^2) &= \sum_{n=1}^N \log p(\mathbf{x}_n \mid \boldsymbol{\mu}, \mathbf{W}, \sigma^2) \\ &= -\frac{ND}{2} \log(2\pi) - \frac{N}{2} \log |\mathbf{C}| - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu})^T \mathbf{C}^{-1} (\mathbf{x}_n - \boldsymbol{\mu}) \end{aligned}$$

Solve for $\boldsymbol{\mu}$

Differentiate the log-likelihood w.r.t. $\boldsymbol{\mu}$ gives

$$\hat{\boldsymbol{\mu}} = \bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

Substituting this back gives the simplified objective in terms of the sample covariance matrix $\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T$:

$$\log p(\mathbf{X} \mid \boldsymbol{\mu}, \mathbf{W}, \sigma^2) = -\frac{N}{2} [D \log(2\pi) + \log |\mathbf{C}| + \text{tr}(\mathbf{C}^{-1} \mathbf{S})]$$

Solve for $\mathbf{W}$ and σ^2 jointly

Rather than iterative fitting, the maximizer is derived directly (Tipping & Bishop 1999):

$$\mathbf{W} = \mathbf{U}_L (\mathbf{L}_L - \sigma^2 \mathbf{I})^{1/2} \mathbf{R}$$

where:

- $\mathbf{U}_L$: $D \times L$ matrix of the top L eigenvectors of $\mathbf{S}$
- $\mathbf{L}_L$: $L \times L$ diagonal matrix of the corresponding top L eigenvalues
- $\mathbf{R}$: arbitrary $L \times L$ rotation matrix (WLOG set $\mathbf{R} = \mathbf{I}$)

and the noise variance:

$$\hat{\sigma}^2 = \frac{1}{D-L} \sum_{i=L+1}^D \lambda_i$$

i.e., the *average of the discarded eigenvalues* (the variance in directions not kept as signal).

Both $\hat{\mathbf{W}}$ and $\hat{\sigma}^2$ come from *one* eigendecomposition of $\mathbf{S}$: compute all D eigenvalues/eigenvectors once, keep the top L for $\hat{\mathbf{W}}$, and average the rest for $\hat{\sigma}^2$.

Sanity check

As $\sigma^2 \rightarrow 0$:

$$\hat{\mathbf{W}} \rightarrow \mathbf{U}_L \mathbf{L}_L^{1/2}$$

which is proportional to the PCA solution.

Getting the latent representation

Unlike PCA's deterministic projection, PPCA computes the *posterior* $p(\mathbf{z} \mid \mathbf{x})$ via Bayes' rule:

$$p(\mathbf{z} \mid \mathbf{x}) = \frac{p(\mathbf{x} \mid \mathbf{z})p(\mathbf{z})}{\int p(\mathbf{x} \mid \mathbf{z}')p(\mathbf{z}') d\mathbf{z}'} = \mathcal{N}(\mathbf{z} \mid \mathbf{M}^{-1}\mathbf{W}^T(\mathbf{x} - \boldsymbol{\mu}), \sigma^2\mathbf{M}^{-1})$$

where $p(\mathbf{z}) = \mathcal{N}(\mathbf{z} \mid \mathbf{0}, \mathbf{I})$, $p(\mathbf{x} \mid \mathbf{z}) = \mathcal{N}(\mathbf{x} \mid \mathbf{W}\mathbf{z} + \boldsymbol{\mu}, \sigma^2\mathbf{I})$ and $\mathbf{M} := \mathbf{W}^T\mathbf{W} + \sigma^2\mathbf{I}$.

The posterior mean serves as the latent embedding (PPCA's analogue of PCA's $\mathbf{z}_n = \mathbf{W}^T\mathbf{x}_n$). In the $\sigma^2 \rightarrow 0$ limit:

$$\mathbb{E}[\mathbf{z} \mid \mathbf{x}] = (\mathbf{W}^T\mathbf{W})^{-1}\mathbf{W}^T(\mathbf{x} - \bar{\mathbf{x}})$$

which reduces exactly to the orthogonal projection used in standard PCA.

8.6 Summary

	PCA	PPCA	Factor Analysis
<i>Decoder</i>	$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n$ (deterministic)	$\mathbf{x} = \mathbf{W}\mathbf{z} + \boldsymbol{\mu} + \boldsymbol{\epsilon}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$	$\mathbf{x} = \mathbf{W}\mathbf{z} + \boldsymbol{\mu} + \boldsymbol{\epsilon}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Psi})$
<i>Noise structure</i>	none	isotropic: same variance σ^2 in every dimension	diagonal: different variance ψ_i per dimension
<i>Prior on $\mathbf{z}$</i>	none	$p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$	$p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$
<i>Constraint on $\mathbf{W}$</i>	orthonormal columns	orthonormal columns	none
<i>Encoder (inference)</i>	deterministic projection: $\mathbf{z}_n = \mathbf{W}^T \mathbf{x}_n$	posterior mean: $\mathbb{E}[\mathbf{z} \mathbf{x}] = \mathbf{M}^{-1} \mathbf{W}^T (\mathbf{x} - \boldsymbol{\mu})$, $\mathbf{M} = \mathbf{W}^T \mathbf{W} + \sigma^2 \mathbf{I}$	posterior mean: $\mathbb{E}[\mathbf{z} \mathbf{x}] = \mathbf{W}^T \boldsymbol{\Sigma}_x^{-1} (\mathbf{x} - \boldsymbol{\mu})$, using full $\boldsymbol{\Sigma}_x = \mathbf{W} \mathbf{W}^T + \boldsymbol{\Psi}$
<i>Encoder uncertainty</i>	none	yes: $\text{Cov}(\mathbf{z} \mathbf{x}) = \sigma^2 \mathbf{M}^{-1}$	yes: full posterior covariance, generally non-diagonal
<i>Fitting method</i>	eigendecomposition of $\mathbf{S}$ (direct, no likelihood)	closed-form MLE via eigendecomposition of $\mathbf{S}$ (Tipping & Bishop)	no closed form — requires EM algorithm
<i>Marginal $p(\mathbf{x})$</i>	none	$\mathcal{N}(\boldsymbol{\mu}, \mathbf{W} \mathbf{W}^T + \sigma^2 \mathbf{I})$	$\mathcal{N}(\boldsymbol{\mu}, \mathbf{W} \mathbf{W}^T + \boldsymbol{\Psi})$
<i>Generative?</i>	No	Yes	Yes

Table 1: Comparison of PCA, PPCA, and Factor Analysis in encoder/decoder terms.

9 Autoencoders

9.1 Deterministic autoencoders

Recall that PCA and FA as learning a linear mapping from $\mathbf{x} \rightarrow \mathbf{z}$, called the **encoder**, f_e , and learning another (linear) mapping $\mathbf{z} \rightarrow \mathbf{x}$, called the **decoder**, f_d . The overall reconstruction function has the form $r(\mathbf{x}) = f_d(f_e(\mathbf{x}))$. The model is trained to minimize $\mathcal{L}(\boldsymbol{\theta}) = \|\mathbf{r}(\mathbf{x}) - \mathbf{x}\|_2^2$, if isotropic generative model's noise variance is assumed. More generally, $\mathcal{L}(\boldsymbol{\theta}) = -\log p(\mathbf{x} | r(\mathbf{x}))$.

Autoencoders uses neural networks to implement non-linear mappings in the encoder and decoder.




Figure 2: An autoencoder with one hidden layer MLP

The hidden units in the middle is a low-dimensional bottleneck between the input and its reconstruction.

If the hidden layer is wide enough, the model may learn the identity function. Thus, a narrow bottleneck layer with $L \ll D$ can be used, called **undercomplete representation**; a wider bottleneck layer with $L \gg D$ can also be used but to impose other kind of regularization, called **overcomplete representation**.

Regularization techniques to avoid trivial memorization,

1. **Bottleneck autoencoders:** This is a special case of a **linear autoencoder**, in which there is one hidden layer, the hidden units are computed using $\mathbf{z} = \mathbf{W}_1 \mathbf{x}$, and the output is reconstructed using $\hat{\mathbf{x}} = \mathbf{W}_2 \mathbf{z}$, where $\mathbf{W}_1$ is a $L \times D$ matrix, $\mathbf{W}_2$ is a $D \times L$ matrix, and $L < D$ (restricting model capacity). Hence $\hat{\mathbf{x}} = \mathbf{W}_2 \mathbf{W}_1 \mathbf{x} = \mathbf{W} \mathbf{x}$ is the output of the model. If we train this model to minimize the squared reconstruction error, $\mathcal{L}(\mathbf{W}) = \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{W} \mathbf{x}_n\|_2^2$, literature shows that $\hat{\mathbf{W}}$ is an orthogonal projection onto the first L eigenvectors of the empirical covariance matrix of the data, equivalent to PCA. If non-linearities are introduced, such as replacing MLP with CNN, the model is strictly more powerful than PCA without significantly changes the model size and training time.




Figure 3: Schematic illustration of an undercomplete autoencoder.

2. **Denoising autoencoders:** Adding noises such as Gaussian noise to the input, the model is able to recover details that are missing in the input, since it has seen similar images before, and can store this information in the parameters of the model. Suppose we train a DAE using Gaussian corruption and squared error reconstruction, i.e., we use $p_c(\tilde{\mathbf{x}} | \mathbf{x}) = \mathcal{N}(\tilde{\mathbf{x}} | \mathbf{x}, \sigma^2 \mathbf{I})$ and $\ell(\mathbf{x}, r(\tilde{\mathbf{x}})) = \|\mathbf{r}(\tilde{\mathbf{x}}) - \mathbf{x}\|_2^2$ for sample $\mathbf{x}$. DAE can learn a vector field, corresponding to the gradient of the log data density $\mathbf{e}(\mathbf{x}) = \mathbf{r}(\tilde{\mathbf{x}}) - \mathbf{x} \approx \nabla_{\mathbf{x}} \log p(\mathbf{x})$. Thus points that are close to the data manifold will be projected onto it via the sampling process.
3. **Contractive autoencoders:** Adding a penalty term to the reconstruction loss, $\lambda \left\| \frac{\partial f_e(\mathbf{x})}{\partial \mathbf{x}} \right\|_F^2 = \lambda \sum_k \|\nabla_{\mathbf{x}} h_k(\mathbf{x})\|_2^2$ (i.e. penalize the Frobenius norm of the encoder's Jacobian), forcing the encoding to be insensitive to small input perturbations.
4. **Sparse autoencoders:** Adding a penalty term to the reconstruction loss, $\lambda \|\mathbf{z}\|_1$, forcing some embeddings to be small/zero. Thus, even under the overcomplete representation ($L > D$), it restricts how much of the model capacity gets used at any one time, which can actually give a richer, more interpretable representation

In summary, the workflow of a deterministic autoencoder is as follows

Setup

- Input $\mathbf{x} \in \mathbb{R}^D$ (high-dim), latent $\mathbf{z} \in \mathbb{R}^K$ with $K \ll D$
- Encoder $f_e(\mathbf{x}; \boldsymbol{\phi}) : \mathbb{R}^D \rightarrow \mathbb{R}^K$
- Decoder $f_d(\mathbf{z}; \boldsymbol{\theta}) : \mathbb{R}^K \rightarrow \mathbb{R}^D$

Training steps

1. Sample a minibatch $\{\mathbf{x}_i\}$ from the dataset
2. Forward pass: $\mathbf{z}_i = f_e(\mathbf{x}_i; \boldsymbol{\phi})$ (deterministic, single vector — no sampling)
3. Reconstruct: $\hat{\mathbf{x}}_i = f_d(\mathbf{z}_i; \boldsymbol{\theta})$

4. Compute loss:

$$\mathcal{L} = \frac{1}{N} \sum_i \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2$$

(or cross-entropy, depending on data type)

5. Backprop $\nabla_{\boldsymbol{\theta}} \mathcal{L}, \nabla_{\boldsymbol{\phi}} \mathcal{L}$, and update $\boldsymbol{\theta}, \boldsymbol{\phi}$ with an optimizer (Adam, SGD, etc.)

6. Repeat over epochs until convergence

Inference Taking a new/held-out point $\mathbf{x}$, passing it through the encoder $\mathbf{z} = f_e(\mathbf{x}; \boldsymbol{\phi})$, where $\mathbf{z} \in \mathbb{R}^K$ is the reduced-dimension representation. Passing through the decoder if reconstruction is needed.

9.2 Variational autoencoders (VAE)

	Deterministic	Probabilistic
<i>Linear</i>	PCA	PPCA, FA
<i>Nonlinear</i>	Bottleneck/Denoising/Contractive/Sparse AE	VAE

Table 2: Comparison of dimensionality reduction algorithms.

VAE can be thought of as a probabilistic version of the deterministic autoencoder. Thus, as a generative model, it can generate new samples instead of just computing embeddings of input vectors.

Also, VAE is the non-linear extension of the factor analysis generative model. Replacing $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{x} | \mathbf{W}\mathbf{z}, \sigma^2 \mathbf{I})$ with

$$p_{\boldsymbol{\theta}}(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{x} | f_d(\mathbf{z}; \boldsymbol{\theta}), \sigma^2 \mathbf{I})$$

where f_d is a non-linear decoder with parameter $\boldsymbol{\theta}$.

Since the posterior $p(\mathbf{z} | \mathbf{x})$ is intractable, so another model $q(\mathbf{z} | \mathbf{x})$ is created to approximate the posterior, called recognition network or inference network. If the posterior is assumed to be Gaussian, with diagonal covariance, then,

$$q_{\boldsymbol{\phi}}(\mathbf{z} | \mathbf{x}) = \mathcal{N}(\mathbf{z} | f_{e,\mu}(\mathbf{x}; \boldsymbol{\phi}), \text{diag}(f_{e,\sigma}(\mathbf{x}; \boldsymbol{\phi})))$$

where $f_{e,\mu/\sigma}$ is the encoder for mean and variance respectively with parameter $\boldsymbol{\phi}$. This idea of training an inference network to invert a generative network, rather than running an optimization algorithm to infer the latent code, is called amortized inference.




Figure 4: Schematic illustration of a VAE.

Training VAEs

The computation of exact $p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \int p_{\theta}(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$ for MLE training is in possible because the porterior in a nonlinear FA model is intractable. However, by making use of $q(\mathbf{z} | \mathbf{x})$, ELBO can be computed. For a single example $\mathbf{x}$, the ELBO as the objective function is,

$$\begin{aligned} \mathcal{L}(\theta, \phi | \mathbf{x}) &= \mathbb{E}_{q_{\phi}(\mathbf{z} | \mathbf{x})} [\log p_{\theta}(\mathbf{x}, \mathbf{z}) - \log q_{\phi}(\mathbf{z} | \mathbf{x})] \\ &= \mathbb{E}_{q(\mathbf{z} | \mathbf{x}, \phi)} [\log p(\mathbf{x} | \mathbf{z}, \theta)] - D_{KL}(q(\mathbf{z} | \mathbf{x}, \phi) \| p(\mathbf{z})) \end{aligned}$$

where $q_{\phi}(\mathbf{z} | \mathbf{x}) = q(\mathbf{z} | \mathbf{x}, \phi)$ are different notation conventions. This can be interpreted as the expected log-likelihood plus a regularizer which penalizes the posterior from deviating too much from the prior.

ELBO is a lower bound of the evidence as shown by the Jensen's inequality,

$$\begin{aligned} \mathcal{L}(\theta, \phi | \mathbf{x}) &= \int q_{\phi}(\mathbf{z} | \mathbf{x}) \log \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z} | \mathbf{x})} d\mathbf{z} \\ &\leq \log \int q_{\phi}(\mathbf{z} | \mathbf{x}) \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z} | \mathbf{x})} d\mathbf{z} = \log p_{\theta}(\mathbf{x}) \end{aligned}$$

Thus, the optimization task for training VAE is

$$\theta^*, \phi^* = \arg \max_{\theta, \phi} \sum_i \mathcal{L}(\theta, \phi | \mathbf{x}_i)$$

A trick in computing ELBO is **reparameterization**. Since $q_{\phi}(\mathbf{z} | \mathbf{x})$ is Gaussian, then

$$\mathbf{z} = f_{e, \mu}(\mathbf{x}; \phi) + f_{e, \sigma}(\mathbf{x}; \phi) \odot \epsilon$$

where $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ and $\odot$ denotes element-wise multiplication. Hence,

$$\mathcal{L}(\theta, \phi | \mathbf{x}) = \mathbb{E}_{\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} [\log p_{\theta}(\mathbf{x} | \mathbf{z} = \mu_{\phi}(\mathbf{x}) + \sigma_{\phi}(\mathbf{x}) \odot \epsilon)] - D_{KL}(q_{\phi}(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z}))$$

The expectation is now over a fixed distribution $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. Before reparameterization, every gradient step would change $q_{\phi}(\mathbf{z} | \mathbf{x})$ itself (the ground samples are taking from shifts). After reparameterization, always sampling from the same fixed source of randomness ϵ , and all the

parameter(ϕ)-dependence has been moved into a deterministic, differentiable path inside the expectation (i.e. once ϵ is sampled, only a deterministic and differentiable neural network $f_{e,\mu/\sigma}$ is left). Therefore, the backpropagation for training can be used as usual.

The first term in the ELBO can be approximated by sampling ϵ , scaling it by the output of the inference network to get z , and then evaluating $\log p(x|z)$ using the decoder network f_d .

The second term in the ELBO is the KL of two Gaussians, which has a closed form expression. With $p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$ and $q_\phi(z|x) = \mathcal{N}(z|\mu, \text{diag}(\sigma))$,

$$D_{KL}(q||p) = \sum_{k=1}^K \left[\log \left(\frac{1}{\sigma_k} \right) + \frac{\sigma_k^2 + (\mu_k - 0)^2}{2 \cdot 1} - \frac{1}{2} \right] = -\frac{1}{2} \sum_{k=1}^K [\log \sigma_k^2 - \sigma_k^2 - \mu_k^2 + 1]$$




Figure 5: Computation graphs for VAEs. Computation graph for VAEs, where $p(z) = \mathcal{N}(z|\mathbf{0}, \mathbf{I})$, $p(x|z, \theta) = \mathcal{N}(x|f(z), \sigma^2 \mathbf{I})$, and $q(z|x, \phi) = \mathcal{N}(z|\mu(x), \Sigma(x))$. Red boxes show sampling operations which are not differentiable. Blue boxes show loss layers (we assume Gaussian likelihoods and priors). (Left) Without the reparameterization trick. (Right) With the reparameterization trick. Gradients can flow from the output loss, back through the decoder and into the encoder.

In summary,

Setup

- $x \in \mathbb{R}^D$, $z \in \mathbb{R}^K$
- Encoder outputs distribution params: $\mu_\phi(x), \sigma_\phi(x) \in \mathbb{R}^K$
- Decoder $f_d(z; \theta)$ outputs likelihood params for $p_\theta(x|z)$
- Fixed prior $p(z) = \mathcal{N}(0, I)$

Training steps

1. Sample a minibatch $\{x_i\}$, then forward pass through encoder: $(\mu_\phi(x_i), \sigma_\phi(x_i)) = f_e(x_i; \phi)$
2. Reparameterize: $z_i = \mu_\phi(x_i) + \sigma_\phi(x_i) \odot \epsilon$ (deterministic given ϵ , differentiable w.r.t. ϕ) with a sample from $\epsilon \sim \mathcal{N}(0, I)$
3. Decode: pass z_i through decoder to get $p_\theta(x|z_i)$'s parameters, which can be used for ELBO first term calculation.
4. Compute ELBO:

$$\mathcal{L}(\theta, \phi | x_i) = \log p_\theta(x_i | z_i) - D_{KL}(q_\phi(z | x_i) || p(z))$$

5. Minimize $-\frac{1}{N} \sum_i \mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\phi} \mid \mathbf{x}_i)$ via backprop — gradients flow through the reparameterized $\mathbf{z}_i$ into both $\boldsymbol{\phi}$ and $\boldsymbol{\theta}$, and update $\boldsymbol{\theta}, \boldsymbol{\phi}$ with an optimizer
6. Repeat over epochs

Inference Since the encoder gives a *distribution*, not a point:

- Option A: use the mean only, $\mathbf{z} = \mu_{\boldsymbol{\phi}}(\mathbf{x})$, as the reduced representation — deterministic, ignores $\sigma_{\boldsymbol{\phi}}(\mathbf{x})$ entirely.
- Option B: sample $\mathbf{z} \sim q_{\boldsymbol{\phi}}(\mathbf{z} \mid \mathbf{x}) = \mathcal{N}(\mu_{\boldsymbol{\phi}}(\mathbf{x}), \sigma_{\boldsymbol{\phi}}(\mathbf{x})^2)$ each time — introduces stochasticity into the reduced representation.

Key property: because of the KL regularization during training, the resulting $\mathbf{z}$ -space (using $\mu_{\boldsymbol{\phi}}(\mathbf{x})$) tends to be smoother and better organized around the origin than a plain AE's latent space, at some cost to reconstruction fidelity (since the KL term competes with the reconstruction term).

Part IV

Kernel Methods

Kernel methods are a class of nonparametric methods for regression and classification. Such methods do not assume a fixed parametric form for the prediction function, but instead try to estimate the function itself directly from data.

For a given dataset $\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^N$, where $y_n = f(\mathbf{x}_n)$ and f is unknown. To predict the function value at a new point $\mathbf{x}_*$, kernel methods predicts that $f(\mathbf{x}_*)$ with some weighted combination of the $\{f(\mathbf{x}_n)\}$ values, i.e. remembering the entire training set $\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^N$ to make predictions.

10 Mercer Kernels and Mercer's Theorem

The key to nonparametric methods is that we need a way to encode prior knowlegde about the similarity of two input vectors. To define similarity, we use a **kernel function**.

Here, consider a **Mercer kernel**, also called a positive definite kernel. This is any symmetric function $\mathcal{K} : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}^+$ such that

$$\sum_{i=1}^N \sum_{j=1}^N c_i c_j \mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) \geq 0$$

for all $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathcal{X}$ which are vectors in $\mathbb{R}^D$ and any choice of numbers $c_i \in \mathbb{R}$. Assuming $\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) > 0$, so that only equality if $c_i = 0$ for all i .

Another way to understand this condition: given a set of N datapoints, define the Gram matrix as the following $N \times N$ similarity matrix:

$$\mathbf{K} = \begin{bmatrix} \mathcal{K}(\mathbf{x}_1, \mathbf{x}_1) & \cdots & \mathcal{K}(\mathbf{x}_1, \mathbf{x}_N) \\ \vdots & \ddots & \vdots \\ \mathcal{K}(\mathbf{x}_N, \mathbf{x}_1) & \cdots & \mathcal{K}(\mathbf{x}_N, \mathbf{x}_N) \end{bmatrix} \in \mathbb{R}^{N \times N}$$

then, the condition becomes that $\mathbf{K}$ is positive definite, i.e. all its eigenvalues are positive, for any set of distinct inputs $\{\mathbf{x}_i\}_{i=1}^N$.

10.1 Mercer's theorem

Any symmetric matrix $\mathbf{K}$ can be represented using an eigendecomposition $\mathbf{K} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$ where $\mathbf{U}$ is an orthonormal matrix of eigenvectors and $\mathbf{\Lambda}$ is a diagonal matrix of eigenvalues $\lambda_i > 0$.

Consider an element (i, j) of $\mathbf{K}$,

$$k_{ij} = \mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{\Lambda}^{\frac{1}{2}} \mathbf{U}_{:,i})^\top (\mathbf{\Lambda}^{\frac{1}{2}} \mathbf{U}_{:,j})$$

where $\mathbf{U}_{:,i}$ is the i -th column of $\mathbf{U}$. Define a mapping $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^N$ as $\phi(\mathbf{x}_i) = \mathbf{\Lambda}^{\frac{1}{2}} \mathbf{U}_{:,i}$. Then,

$$k_{ij} = \mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j) = \sum_{m=1}^N \phi_m(\mathbf{x}_i) \phi_m(\mathbf{x}_j)$$

In theory, the kernel function can be generalized to $N \rightarrow \infty$ with $\lambda_i > 0$ as a prerequisite.

ϕ is a mapping from input space to the feature space, which is a Hilbert space. **The kernel function $\mathcal{K}$ can be interpreted as an inner product in the feature space**, i.e. $\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle$. The kernel trick is that ϕ is never explicitly computed, but from eigendecomposition of $\mathbf{K}$.

10.2 Mercer kernels examples

For real-valued inputs $\mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^D$, it is common to use stationary kernels, which are functions of the form $\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \mathcal{K}(\|\mathbf{x}_i - \mathbf{x}_j\|)$, i.e. the value depends only on the elementwise distance between the two inputs.

Gaussian Kernel

Also called the squared exponential kernel, exponentiated quadratic kernel, RBF kernel, defined as

$$\mathcal{K}(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right)$$

ℓ is the length scale of the kernel, i.e., the distance over which we expect differences to matter. This is known as the bandwidth parameter.

ARD Kernel

Automatic relevancy determination kernel is a generalization of the Gaussian kernel by replacing Euclidean distance with Mahalanobis distance, defined as

$$\mathcal{K}(\mathbf{x}, \mathbf{x}') = \sigma^2 \exp\left(-\frac{1}{2} \mathbf{r}^\top \boldsymbol{\Sigma}^{-1} \mathbf{r}\right)$$

Typically, $\boldsymbol{\Sigma}$ is a diagonal matrix with ℓ_i^2 on the diagonal, thus

$$\mathcal{K}(\mathbf{x}, \mathbf{x}') = \sigma^2 \exp\left(-\sum_{i=1}^D \frac{(x_i - x'_i)^2}{2\ell_i^2}\right)$$

which allows each input dimension i to have its own characteristic length scale. For high-dimensional inputs where some dimensions may be more relevant than others (eg. $\ell_i = \infty$ makes that dimension been ignored).

10.3 Combining kernels

Multiplying two positive-definite kernels together always results in another positive definite kernel. This is a way to get a conjunction of the individual properties of each kernel.

$$\mathcal{K}(\mathbf{x}, \mathbf{x}') = \mathcal{K}_1(\mathbf{x}, \mathbf{x}') \times \mathcal{K}_2(\mathbf{x}, \mathbf{x}')$$

Adding two positive-definite kernels together always results in another positive definite kernel. This is a way to get a disjunction of the individual properties of each kernel.

$$\mathcal{K}(\mathbf{x}, \mathbf{x}') = \mathcal{K}_1(\mathbf{x}, \mathbf{x}') + \mathcal{K}_2(\mathbf{x}, \mathbf{x}')$$

11 Gaussian Processes

Gaussian processes is a way to define distributions over functions of the form $f : \mathcal{X} \rightarrow \mathbb{R}$, where $\mathcal{X}$ is any domain. The key assumption is that the function values at a set of $M > 0$ inputs, $\mathbf{f} = [f(\mathbf{x}_1), \dots, f(\mathbf{x}_M)]$, is jointly Gaussian, with mean $(\boldsymbol{\mu} = m(\mathbf{x}_1), \dots, m(\mathbf{x}_M))$ and covariance $\Sigma_{ij} = \mathcal{K}(\mathbf{x}_i, \mathbf{x}_j)$, where m is a mean function and $\mathcal{K}$ is a positive definite (Mercer) kernel.

Since we assume this holds for any $M > 0$, this includes the case where $M = N + 1$, containing N training points $\mathbf{x}_n$ and 1 test point $\mathbf{x}_*$. Then, $f(\mathbf{x}_*)$ can be inferred from knowledge of $f(\mathbf{x}_1), \dots, f(\mathbf{x}_N)$ by manipulating the joint Gaussian distribution $p(f(\mathbf{x}_1), \dots, f(\mathbf{x}_N), f(\mathbf{x}_*))$. Also, extending this to noisy functions of $f(\mathbf{x}_n)$ can deal with problems such as regression or classification.




Figure 6: A Gaussian process for 2 training points and 1 test point. The hidden nodes $f_i = f(\mathbf{x}_i)$ represent the value of the function at each of the data points. These hidden nodes are fully interconnected with the edge represent the covariance terms $\Sigma_{ij} = \mathcal{K}(\mathbf{x}_i, \mathbf{x}_j)$. If the testing points $\mathbf{x}_*$ is close to the training points, then the predicted function value $f(\mathbf{x}_*)$ will be similar to the training points $f(\mathbf{x}_1)$ and $f(\mathbf{x}_2)$.

11.1 Noise-free observations

Suppose we observe a training set $\mathcal{D} = \{(\mathbf{x}_n, y_n) : n = 1 : N\}$, where $y_n = f(\mathbf{x}_n)$ is the noise-free observation of the function evaluated at $\mathbf{x}_n$. If we ask the GP to predict $f(\mathbf{x})$ for a value of $\mathbf{x}$ that it has already seen, we want the GP to return the answer $f(\mathbf{x})$ with no uncertainty.

Now for the case of predicting the outputs for new inputs. Specifically, given a test set $\mathbf{X}_*$ with N_* new data points, we want to predict the function outputs $\mathbf{f}_* = [f(\mathbf{x}_{*1}), \dots, f(\mathbf{x}_{*N_*})]$. By definition of the GP, the joint distribution $p(\mathbf{f}_X, \mathbf{f}_* | \mathbf{X}, \mathbf{X}_*)$ has the following form

$$\begin{pmatrix} \mathbf{f}_X \\ \mathbf{f}_* \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} \boldsymbol{\mu}_X \\ \boldsymbol{\mu}_* \end{pmatrix}, \begin{pmatrix} \mathbf{K}_{X,X} & \mathbf{K}_{X,*} \\ \mathbf{K}_{X,*}^\top & \mathbf{K}_{*,*} \end{pmatrix} \right)$$

where $\boldsymbol{\mu}_X = [m(\mathbf{x}_1), \dots, m(\mathbf{x}_N)]$, $\boldsymbol{\mu}_* = [m(\mathbf{x}_{*1}), \dots, m(\mathbf{x}_{*N_*})]$, $\mathbf{K}_{X,X} = \mathcal{K}(\mathbf{X}, \mathbf{X})$ is $N \times N$, $\mathbf{K}_{X,*} = \mathcal{K}(\mathbf{X}, \mathbf{X}_*)$ is $N \times N_*$, and $\mathbf{K}_{*,*} = \mathcal{K}(\mathbf{X}_*, \mathbf{X}_*)$ is $N_* \times N_*$. By the standard rules for conditioning Gaussians, the posterior has the following form

$$\begin{aligned} p(\mathbf{f}_* | \mathbf{X}_*, \mathcal{D}) &= \mathcal{N}(\mathbf{f}_* | \boldsymbol{\mu}_{*|X}, \boldsymbol{\Sigma}_{*|X}) \\ \boldsymbol{\mu}_{*|X} &= \boldsymbol{\mu}_* + \mathbf{K}_{X,*}^\top \mathbf{K}_{X,X}^{-1} (\mathbf{f}_X - \boldsymbol{\mu}_X) \\ \boldsymbol{\Sigma}_{*|X} &= \mathbf{K}_{*,*} - \mathbf{K}_{X,*}^\top \mathbf{K}_{X,X}^{-1} \mathbf{K}_{X,*} \end{aligned}$$

11.2 Noisy observations

Suppose we observe a training set $\mathcal{D} = \{(\mathbf{x}_n, y_n) : n = 1 : N\}$, where $y_n = f(\mathbf{x}_n) + \epsilon_n$ is the noisy observation of the function evaluated at $\mathbf{x}_n$, and $\epsilon_n \sim \mathcal{N}(0, \sigma_y^2)$ is i.i.d. Gaussian noise. Consider, the covariance of the observed noisy responses is

$$\begin{aligned}\text{Cov}[y_i, y_j] &= \text{Cov}[f(\mathbf{x}_i) + \epsilon_i, f(\mathbf{x}_j) + \epsilon_j] \\ &= \text{Cov}[f(\mathbf{x}_i), f(\mathbf{x}_j)] + \text{Cov}[\epsilon_i, \epsilon_j] \\ &= \mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) + \sigma_y^2 \delta_{ij}\end{aligned}$$

Writing them compactly as,

$$\text{Cov}[\mathbf{y} \mid \mathbf{X}] = \mathbf{K}_{X,X} + \sigma_y^2 \mathbf{I}_N \triangleq \mathbf{K}_\sigma$$

The joint density of the observed data and the latent function on the test points is given by

$$\begin{pmatrix} \mathbf{y} \\ \mathbf{f}_* \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} \boldsymbol{\mu}_X \\ \boldsymbol{\mu}_* \end{pmatrix}, \begin{pmatrix} \mathbf{K}_\sigma & \mathbf{K}_{X,*} \\ \mathbf{K}_{X,*}^\top & \mathbf{K}_{*,*} \end{pmatrix} \right)$$

Therefore, on a set of test points $\mathbf{X}_*$, the posterior predictive distribution is

$$\begin{aligned}p(\mathbf{f}_* | \mathbf{X}_*, \mathcal{D}) &= \mathcal{N}(\mathbf{f}_* | \boldsymbol{\mu}_{*,X}, \boldsymbol{\Sigma}_{*,X}) \\ \boldsymbol{\mu}_{*,X} &= \boldsymbol{\mu}_* + \mathbf{K}_{X,*}^\top \mathbf{K}_\sigma^{-1} (\mathbf{f}_X - \boldsymbol{\mu}_X) \\ \boldsymbol{\Sigma}_{*,X} &= \mathbf{K}_{*,*} - \mathbf{K}_{X,*}^\top \mathbf{K}_\sigma^{-1} \mathbf{K}_{X,*}\end{aligned}$$

11.3 Bayesian linear regression

Bayesian linear regression is a special case of Gaussian processes, and can be viewed in weight space or function space. In the following example, let's consider the case where we have N training points $\mathcal{D} = \{(\mathbf{x}_n, \mathbf{y}_n) : n = 1 : N\}$ and 1 testing points $\mathbf{x}_*$.

Weight space view

Consider the linear regression model $y = f(\mathbf{x}) + \epsilon$, where $f(\mathbf{x}) = \mathbf{w}^\top \boldsymbol{\phi}(\mathbf{x})$ and $\epsilon \sim \mathcal{N}(0, \sigma_y^2)$. With vectorized form:

$$\mathbf{y} = \boldsymbol{\Phi} \mathbf{w} + \boldsymbol{\epsilon} \implies p(\mathbf{y} | \mathbf{w}) = \mathcal{N}(\mathbf{y} | \boldsymbol{\Phi} \mathbf{w}, \sigma_y^2 \mathbf{I})$$

where $\boldsymbol{\Phi}$ is the $N \times D$ design matrix including features of all N training points. Plain linear regression stops by maximizing this w.r.t. $\mathbf{w}$, but the Bayesian pushes this forward.

Assuming a Gaussian prior $p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \boldsymbol{\Sigma}_w)$ on the model parameter $\mathbf{w} \in \mathbb{R}^{D \times 1}$, then the posterior is as follows:

$$p(\mathbf{w} | \mathcal{D}) = \mathcal{N} \left(\mathbf{w} \mid \frac{1}{\sigma_y^2} \mathbf{A}^{-1} \boldsymbol{\Phi}^\top \mathbf{y}, \mathbf{A}^{-1} \right) \propto p(\mathbf{y} | \mathbf{w}) p(\mathbf{w})$$

which is derived by completing the square with $\mathbf{A} = \sigma_y^{-2} \boldsymbol{\Phi}^\top \boldsymbol{\Phi} + \boldsymbol{\Sigma}_w^{-1}$.

The posterior predictive distribution for $f_* = f(\mathbf{x}_*) = \boldsymbol{\phi}_*^\top \mathbf{w} = \boldsymbol{\phi}(\mathbf{x}_*)^\top \mathbf{w}$ is therefore

$$p(f_* | \mathcal{D}, \mathbf{x}_*) = \mathcal{N} \left(f_* \mid \frac{1}{\sigma_y^2} \boldsymbol{\phi}_*^\top \mathbf{A}^{-1} \boldsymbol{\Phi}^\top \mathbf{y}, \boldsymbol{\phi}_*^\top \mathbf{A}^{-1} \boldsymbol{\phi}_* \right)$$

using the standard rules for linear transforms of Gaussians ($\mathbf{z} = \mathbf{c}^\top \mathbf{w} \implies \mathbb{E}[\mathbf{z}] = \mathbf{c}^\top \boldsymbol{\mu}_w$, $\text{Cov}(\mathbf{z}) = \mathbf{c}^\top \boldsymbol{\Sigma}_w \mathbf{c}$).

Function space view

Bayesian linear regression can be viewed as a GP with a kernel $\mathcal{K}(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^\top \Sigma_w \phi(\mathbf{x}')$, which measures the similarity between two inputs $\mathbf{x}$ and $\mathbf{x}'$ in the feature space defined by ϕ and is also a valid kernel since covariance Σ_w is always positive semi-definite.

Then, defining the training-training, training-test, and test-test covariance:

$$\mathbf{K} = \Phi^\top \Sigma_w \Phi \in \mathbb{R}^{N \times N}, \quad \mathbf{k}_* = \Phi^\top \Sigma_w \phi_* \in \mathbb{R}^{N \times 1}, \quad k_{**} = \phi_*^\top \Sigma_w \phi_* \in \mathbb{R}$$

where ϕ_* is a $D \times 1$ is the feature vector of the single test point. Using these notation, and the matrix inversion lemma, we can rewrite the predictive posterior in the weight space view section as:

$$\begin{aligned} p(f_* | \mathcal{D}, \mathbf{x}_*) &= \mathcal{N}(f_* | \mu_{*|X}, \Sigma_{*|X}) \\ \mu_{*|X} &= \phi_*^\top \Sigma_w \Phi^\top (\mathbf{K} + \sigma_y^2 \mathbf{I})^{-1} \mathbf{y} = \mathbf{k}_*^\top \mathbf{K}_\sigma^{-1} \mathbf{y} \\ \Sigma_{*|X} &= \phi_*^\top \Sigma_w \phi_* - \phi_*^\top \Sigma_w \Phi^\top (\mathbf{K} + \sigma_y^2 \mathbf{I})^{-1} \Phi \Sigma_w \phi_* = k_{**} - \mathbf{k}_*^\top \mathbf{K}_\sigma^{-1} \mathbf{k}_* \end{aligned}$$

which matches the predictive posterior in the GP with noisy observations section with only one testing point $\mathbf{x}_*$ instead of N_* testing points $\mathbf{X}_*$.

With kernel trick and a valid kernel $\mathcal{K}$, we can compute the predictive posterior without explicitly computing the feature mapping ϕ and allowing for a infinite length feature space.

11.4 Estimating the kernel

Most kernels have some free parameters, which can have a large effect on the predictions from the model. To estimate the kernel parameters θ (i.e. hyperparameters), we could use exhaustive search over a discrete grid of values, with validation loss as an objective, but this can be quite slow.

Empirical Bayes

Consider an empirical Bayes approach, which allows the use of gradient-based optimization methods, which are much faster. In particular, we will maximize the marginal likelihood

$$p(\mathbf{y} | \mathbf{X}, \theta) = \int p(\mathbf{y} | \mathbf{f}, \mathbf{X}) p(\mathbf{f} | \mathbf{X}, \theta) d\mathbf{f}$$

The latent Gaussian vector $\mathbf{f}$ have been marginalized out, so this is called the marginal likelihood rather than likelihood.

Recall that the prior $p(\mathbf{f} | \mathbf{X}) = \mathcal{N}(\mathbf{f} | \mathbf{0}, \mathbf{K}_{X,X})$ and the likelihood $p(\mathbf{y} | \mathbf{f}) = \mathcal{N}(\mathbf{y} | \mathbf{f}, \sigma_y^2 \mathbf{I}_N)$, in which we assume the mean function is 0, then the marginal likelihood is given by

$$\log p(\mathbf{y} | \mathbf{X}, \theta) = \log \mathcal{N}(\mathbf{y} | \mathbf{0}, \mathbf{K}_\sigma) = -\frac{1}{2} \mathbf{y}^\top \mathbf{K}_\sigma^{-1} \mathbf{y} - \frac{1}{2} \log |\mathbf{K}_\sigma| - \frac{N}{2} \log(2\pi)$$

where the dependence of $\mathbf{K}_\sigma = \mathbf{K}_{X,X} + \sigma_y^2 \mathbf{I}_N$ on θ is implicit. The first term is a data fit term, the second term is a model complexity term, and the third term is a constant. There is a tradeoff between the first and the second term.

To maximize the marginal likelihood,

$$\begin{aligned} \frac{\partial}{\partial \theta_j} \log p(\mathbf{y} | \mathbf{X}, \theta) &= \frac{1}{2} \mathbf{y}^\top \mathbf{K}_\sigma^{-1} \frac{\partial \mathbf{K}_\sigma}{\partial \theta_j} \mathbf{K}_\sigma^{-1} \mathbf{y} - \frac{1}{2} \text{tr} \left(\mathbf{K}_\sigma^{-1} \frac{\partial \mathbf{K}_\sigma}{\partial \theta_j} \right) \\ &= \frac{1}{2} \text{tr} \left((\alpha \alpha^\top - \mathbf{K}_\sigma^{-1}) \frac{\partial \mathbf{K}_\sigma}{\partial \theta_j} \right) \end{aligned}$$

where $\boldsymbol{\alpha} = \mathbf{K}_\sigma^{-1} \mathbf{y}$. It takes $O(N^3)$ time to compute $\mathbf{K}_\sigma^{-1}$, and then $O(N^2)$ time per hyper-parameter to compute the gradient. Therefore, we can estimate the kernel parameters using any standard gradient-based optimizer. However, since the objective is not convex, local minima can be a problem, so multiple restarts might be used.